

```

import numpy as np
import pdb

"""
This code was based off of code from cs231n at Stanford University,
and modified for ECE C147/C247 at UCLA.
"""
class SVM(object):

    def __init__(self, dims=[10, 3073]):
        self.init_weights(dims=dims)

    def init_weights(self, dims):
        """
        Initializes the weight matrix of the SVM. Note that it has shape
        (C, D)
        where C is the number of classes and D is the feature size.
        """
        self.W = np.random.normal(size=dims)

    def loss(self, X, y):
        """
        Calculates the SVM loss.

        Inputs have dimension D, there are C classes, and we operate on
        minibatches
        of N examples.

        Inputs:
        - X: A numpy array of shape (N, D) containing a minibatch of data.
        - y: A numpy array of shape (N,) containing training labels; y[i]
        = c means
            that X[i] has label c, where 0 <= c < C.

        Returns a tuple of:
        - loss as single float
        """

        # compute the loss and the gradient
        num_classes = self.W.shape[0]
        num_train = X.shape[0]
        loss = 0.0

        for i in np.arange(num_train):
            # =====
            # YOUR CODE HERE:
            #   Calculate the normalized SVM loss, and store it as 'loss'.
            #   (That is, calculate the sum of the losses of all the training
            #   set margins, and then normalize the loss by the number of

```

```

        # training examples.)
# =====
#
    loss_array = 1 + X[i].dot(self.W.T) - X[i].dot(self.W.T)[y[i]]
    loss_array[loss_array < 0] = 0
    loss += np.sum(loss_array, axis=0) - 1 # subtract one to get
rid of the j = y(i) case
    loss /= num_train
# =====
#
# END YOUR CODE HERE
# =====
#

return loss

def loss_and_grad(self, X, y):
    """
    Same as self.loss(X, y), except that it also returns the gradient.

    Output: grad -- a matrix of the same dimensions as W containing
           the gradient of the loss with respect to W.
    """

    # compute the loss and the gradient
    num_classes = self.W.shape[0]
    num_train = X.shape[0]
    loss = 0.0
    grad = np.zeros_like(self.W)

    # =====
#
# YOUR CODE HERE:
# Calculate the SVM loss and the gradient. Store the gradient
in # the variable grad.
# =====
#
for i in np.arange(num_train):
    a = X[i].dot(self.W.T)
    for j in range(num_classes):
        if j == y[i]:
            continue
        zj = 1 + a[j] - a[y[i]]
        if zj > 0:
            loss += zj
            grad[j] += X[i]
            grad[y[i]] -= X[i]

```

```

# =====
#
# END YOUR CODE HERE
# =====
#

    loss /= num_train
    grad /= num_train

    return loss, grad

def grad_check_sparse(self, X, y, your_grad, num_checks=10, h=1e-5):
    """
    sample a few random elements and only return numerical
    in these dimensions.
    """

    for i in np.arange(num_checks):
        ix = tuple([np.random.randint(m) for m in self.W.shape])

        oldval = self.W[ix]
        self.W[ix] = oldval + h # increment by h
        fxph = self.loss(X, y)
        self.W[ix] = oldval - h # decrement by h
        fxmh = self.loss(X,y) # evaluate f(x - h)
        self.W[ix] = oldval # reset

        grad_numerical = (fxph - fxmh) / (2 * h)
        grad_analytic = your_grad[ix]
        rel_error = abs(grad_numerical - grad_analytic) /
        (abs(grad_numerical) + abs(grad_analytic))
        print('numerical: %f analytic: %f, relative error: %e' %
              (grad_numerical, grad_analytic, rel_error))

def fast_loss_and_grad(self, X, y):
    """
    A vectorized implementation of loss_and_grad. It shares the same
    inputs and outputs as loss_and_grad.
    """
    loss = 0.0
    grad = np.zeros(self.W.shape) # initialize the gradient as zero

# =====
#
# YOUR CODE HERE:
# Calculate the SVM loss WITHOUT any for loops.
# =====
#

    # for i in np.arange(num_train):
    #     a = X[i].dot(self.W.T)

```

```

#     for j in range(num_classes):
#         if j == y[i]:
#             continue
#         zj = 1 + a[j] - a[y[i]]
#         if zj > 0:
#             loss += zj
#             grad[j] += X[i]
#             grad[y[i]] -= X[i]

num_train = X.shape[0]

scores = X.dot(self.W.T)
zj = (scores.T - scores[np.arange(num_train), y] + 1).T
zj[zj < 0] = 0
zj[np.arange(num_train), y] = 0
loss = np.sum(zj) / num_train

# =====
#
# END YOUR CODE HERE
# =====
#

# =====
#
# YOUR CODE HERE:
# Calculate the SVM grad WITHOUT any for loops.
# =====
#

mask = np.zeros(zj.shape)
mask[zj > 0] = 1
mask[np.arange(num_train), y] = -np.sum(mask, axis=1)

grad = mask.T.dot(X) / num_train

# =====
#
# END YOUR CODE HERE
# =====
#

return loss, grad

def train(self, X, y, learning_rate=1e-3, num_iters=100,
          batch_size=200, verbose=False):
    """

```

Train this linear classifier using stochastic gradient descent.

Inputs:

- X: A numpy array of shape (N, D) containing training data; there are N training samples each of dimension D.
- y: A numpy array of shape (N,) containing training labels; $y[i] = c$ means that $X[i]$ has label $0 \leq c < C$ for C classes.
- learning_rate: (float) learning rate for optimization.
- num_iters: (integer) number of steps to take when optimizing
- batch_size: (integer) number of training examples to use at each step.
- verbose: (boolean) If true, print progress during optimization.

Outputs:

A list containing the value of the loss function at each training iteration.

```
"""
    num_train, dim = X.shape
    num_classes = np.max(y) + 1 # assume y takes values 0...K-1 where
K is number of classes

    self.init_weights(dims=[np.max(y) + 1, X.shape[1]]) #
initializes the weights of self.W

    # Run stochastic gradient descent to optimize W
    loss_history = []

    for it in np.arange(num_iters):
        X_batch = None
        y_batch = None

        #
===== #
        # YOUR CODE HERE:
        # Sample batch_size elements from the training data for use in
        # gradient descent. After sampling,
        # - X_batch should have shape: (dim, batch_size)
        # - y_batch should have shape: (batch_size,)
        # The indices should be randomly generated to reduce
correlations
        # in the dataset. Use np.random.choice. It's okay to
sample with
        # replacement.
        #
===== #

        indices = np.random.choice(num_train, batch_size)
        X_batch = X[indices]
```

```

        y_batch = y[indices]

        #
        ===== #
        # END YOUR CODE HERE
        #
        ===== #

        # evaluate loss and gradient
        loss, grad = self.fast_loss_and_grad(X_batch, y_batch)
        loss_history.append(loss)

        #
        ===== #
        # YOUR CODE HERE:
        #   Update the parameters, self.W, with a gradient step
        #
        ===== #

        self.W -= learning_rate * grad

        #
        ===== #
        # END YOUR CODE HERE
        #
        ===== #

        if verbose and it % 100 == 0:
            print('iteration {} / {}: loss {}'.format(it, num_iters,
loss))

        return loss_history

def predict(self, X):
    """
    Inputs:
    - X: N x D array of training data. Each row is a D-dimensional
point.

    Returns:
    - y_pred: Predicted labels for the data in X. y_pred is a 1-
dimensional
        array of length N, and each element is an integer giving the
predicted
        class.
    """
    y_pred = np.zeros(X.shape[1])

    # =====

```

```
#
    # YOUR CODE HERE:
    #   Predict the labels given the training data with the parameter
self.W.
    # =====
#
    y_pred = np.argmax(X.dot(self.W.T), axis=1)
    # =====
#
    # END YOUR CODE HERE
    # =====
#

    return y_pred
```