

```

1 import numpy as np
2
3 from .layers import *
4 from .layer_utils import *
5
6 """
7 This code was originally written for CS 231n at Stanford University
8 (cs231n.stanford.edu). It has been modified in various areas for use in the
9 ECE 239AS class at UCLA. This includes the descriptions of what code to
10 implement as well as some slight potential changes in variable names to be
11 consistent with class nomenclature. We thank Justin Johnson & Serena Yeung
12 for
13 permission to use this code. To see the original version, please visit
14 cs231n.stanford.edu.
15 """
16 class TwoLayerNet(object):
17     """
18     A two-layer fully-connected neural network with ReLU nonlinearity and
19     softmax loss that uses a modular layer design. We assume an input dimension
20     of D, a hidden dimension of H, and perform classification over C classes.
21
22     The architecture should be affine - relu - affine - softmax.
23
24     Note that this class does not implement gradient descent; instead, it
25     will interact with a separate Solver object that is responsible for running
26     optimization.
27
28     The learnable parameters of the model are stored in the dictionary
29     self.params that maps parameter names to numpy arrays.
30     """
31
32     def __init__(self, input_dim=3*32*32, hidden_dims=100, num_classes=10,
33                 dropout=0, weight_scale=1e-3, reg=0.0):
34         """
35         Initialize a new network.
36
37         Inputs:
38         - input_dim: An integer giving the size of the input
39         - hidden_dims: An integer giving the size of the hidden layer
40         - num_classes: An integer giving the number of classes to classify
41         - dropout: Scalar between 0 and 1 giving dropout strength.
42         - weight_scale: Scalar giving the standard deviation for random
43           initialization of the weights.
44         - reg: Scalar giving L2 regularization strength.
45         """
46         self.params = {}
47         self.reg = reg
48
49         # ===== #
50         # YOUR CODE HERE:
51         # Initialize W1, W2, b1, and b2. Store these as self.params['W1'],
52         # self.params['W2'], self.params['b1'] and self.params['b2']. The
53         # biases are initialized to zero and the weights are initialized
54         # so that each parameter has mean 0 and standard deviation
55         weight_scale.
56         # The dimensions of W1 should be (input_dim, hidden_dim) and the
57         # dimensions of W2 should be (hidden_dims, num_classes)
58         # ===== #

```

```

59     self.params['W1'] = weight_scale * np.random.randn(input_dim,
hidden_dims)
60     self.params['b1'] = np.zeros(hidden_dims)
61     self.params['W2'] = weight_scale * np.random.randn(hidden_dims,
num_classes)
62     self.params['b2'] = np.zeros(num_classes)
63
64     # ===== #
65     # END YOUR CODE HERE
66     # ===== #
67
68 def loss(self, X, y=None):
69     """
70     Compute loss and gradient for a minibatch of data.
71
72     Inputs:
73     - X: Array of input data of shape (N, d_1, ..., d_k)
74     - y: Array of labels, of shape (N,). y[i] gives the label for X[i].
75
76     Returns:
77     If y is None, then run a test-time forward pass of the model and return:
78     - scores: Array of shape (N, C) giving classification scores, where
79       scores[i, c] is the classification score for X[i] and class c.
80
81     If y is not None, then run a training-time forward and backward pass and
82     return a tuple of:
83     - loss: Scalar value giving the loss
84     - grads: Dictionary with the same keys as self.params, mapping parameter
85       names to gradients of the loss with respect to those parameters.
86     """
87     scores = None
88
89     # ===== #
90     # YOUR CODE HERE:
91     #   Implement the forward pass of the two-layer neural network. Store
92     #   the class scores as the variable 'scores'. Be sure to use the layers
93     #   you prior implemented.
94     # ===== #
95
96     out, l1_cache = affine_relu_forward(X, self.params['W1'],
self.params['b1'])
97     scores, l2_cache = affine_forward(out, self.params['W2'],
self.params['b2'])
98
99     # ===== #
100    # END YOUR CODE HERE
101    # ===== #
102
103    # If y is None then we are in test mode so just return scores
104    if y is None:
105        return scores
106
107    loss, grads = 0, {}
108    # ===== #
109    # YOUR CODE HERE:
110    #   Implement the backward pass of the two-layer neural net. Store
111    #   the loss as the variable 'loss' and store the gradients in the
112    #   'grads' dictionary. For the grads dictionary, grads['W1'] holds
113    #   the gradient for W1, grads['b1'] holds the gradient for b1, etc.
114    #   i.e., grads[k] holds the gradient for self.params[k].

```

```

115 #
116 # Add L2 regularization, where there is an added cost  $0.5 * \text{self.reg} * W^2$ 
117 # for each W. Be sure to include the 0.5 multiplying factor to
118 # match our implementation.
119 #
120 # And be sure to use the layers you prior implemented.
121 # ===== #
122
123 loss, dL = softmax_loss(scores, y)
124 reg_loss = 0.5 * self.reg * (np.sum(self.params['W1']**2) +
np.sum(self.params['W2']**2))
125 loss += reg_loss
126
127 dh1, grads['W2'], grads['b2'] = affine_backward(dL, l2_cache)
128 dx, grads['W1'], grads['b1'] = affine_relu_backward(dh1, l1_cache)
129
130 grads['W1'] += self.reg * self.params['W1']
131 grads['W2'] += self.reg * self.params['W2']
132
133 # ===== #
134 # END YOUR CODE HERE
135 # ===== #
136
137 return loss, grads
138
139
140 class FullyConnectedNet(object):
141     """
142     A fully-connected neural network with an arbitrary number of hidden layers,
143     ReLU nonlinearities, and a softmax loss function. This will also implement
144     dropout and batch normalization as options. For a network with L layers,
145     the architecture will be
146
147     {affine - [batch norm] - relu - [dropout]} x (L - 1) - affine - softmax
148
149     where batch normalization and dropout are optional, and the {...} block is
150     repeated L - 1 times.
151
152     Similar to the TwoLayerNet above, learnable parameters are stored in the
153     self.params dictionary and will be learned using the Solver class.
154     """
155
156     def __init__(self, hidden_dims, input_dim=3*32*32, num_classes=10,
157                   dropout=0, use_batchnorm=False, reg=0.0,
158                   weight_scale=1e-2, dtype=np.float32, seed=None):
159         """
160         Initialize a new FullyConnectedNet.
161
162         Inputs:
163         - hidden_dims: A list of integers giving the size of each hidden layer.
164         - input_dim: An integer giving the size of the input.
165         - num_classes: An integer giving the number of classes to classify.
166         - dropout: Scalar between 0 and 1 giving dropout strength. If dropout=0
then
167     the network should not use dropout at all.
168     - use_batchnorm: Whether or not the network should use batch
normalization.
169     - reg: Scalar giving L2 regularization strength.
170     - weight_scale: Scalar giving the standard deviation for random
171     initialization of the weights.

```

```

172     - dtype: A numpy datatype object; all computations will be performed
using
173     this datatype. float32 is faster but less accurate, so you should use
174     float64 for numeric gradient checking.
175     - seed: If not None, then pass this random seed to the dropout layers.
This
176     will make the dropout layers deterministic so we can gradient check the
177     model.
178     """
179     self.use_batchnorm = use_batchnorm
180     self.use_dropout = dropout > 0
181     self.reg = reg
182     self.num_layers = 1 + len(hidden_dims)
183     self.dtype = dtype
184     self.params = {}
185
186     # ===== #
187     # YOUR CODE HERE:
188     # Initialize all parameters of the network in the self.params
dictionary.
189     # The weights and biases of layer 1 are W1 and b1; and in general the
190     # weights and biases of layer i are Wi and bi. The
191     # biases are initialized to zero and the weights are initialized
192     # so that each parameter has mean 0 and standard deviation
weight_scale.
193     # ===== #
194
195     dims = hidden_dims
196     dims.insert(0, input_dim)
197     dims.append(num_classes)
198
199     for i in range(len(dims) - 1):
200         self.params['W{}'.format(i+1)] = weight_scale *
np.random.randn(dims[i], dims[i+1])
201         self.params['b{}'.format(i+1)] = np.zeros(dims[i+1])
202
203     # ===== #
204     # END YOUR CODE HERE
205     # ===== #
206
207     # When using dropout we need to pass a dropout_param dictionary to each
208     # dropout layer so that the layer knows the dropout probability and the
mode
209     # (train / test). You can pass the same dropout_param to each dropout
layer.
210     self.dropout_param = {}
211     if self.use_dropout:
212         self.dropout_param = {'mode': 'train', 'p': dropout}
213         if seed is not None:
214             self.dropout_param['seed'] = seed
215
216     # With batch normalization we need to keep track of running means and
217     # variances, so we need to pass a special bn_param object to each batch
218     # normalization layer. You should pass self.bn_params[0] to the forward
pass
219     # of the first batch normalization layer, self.bn_params[1] to the
forward
220     # pass of the second batch normalization layer, etc.
221     self.bn_params = []
222     if self.use_batchnorm:

```

```

223     self.bn_params = [{'mode': 'train'} for i in np.arange(self.num_layers
- 1)]
224
225     # Cast all parameters to the correct datatype
226     for k, v in self.params.items():
227         self.params[k] = v.astype(dtype)
228
229
230 def loss(self, X, y=None):
231     """
232     Compute loss and gradient for the fully-connected net.
233
234     Input / output: Same as TwoLayerNet above.
235     """
236     X = X.astype(self.dtype)
237     mode = 'test' if y is None else 'train'
238
239     # Set train/test mode for batchnorm params and dropout param since they
240     # behave differently during training and testing.
241     if self.dropout_param is not None:
242         self.dropout_param['mode'] = mode
243     if self.use_batchnorm:
244         for bn_param in self.bn_params:
245             bn_param[mode] = mode
246
247     scores = None
248
249     # ===== #
250     # YOUR CODE HERE:
251     # Implement the forward pass of the FC net and store the output
252     # scores as the variable "scores".
253     # ===== #
254
255     caches = []
256     x = X
257     for i in range(1, self.num_layers + 1):
258         w, b = self.params['W{}'.format(i)], self.params['b{}'.format(i)]
259         if (i == self.num_layers):
260             scores, cache = affine_forward(x, w, b)
261         else:
262             x, cache = affine_relu_forward(x, w, b)
263         caches.append(cache)
264
265     # ===== #
266     # END YOUR CODE HERE
267     # ===== #
268
269     # If test mode return early
270     if mode == 'test':
271         return scores
272
273     loss, grads = 0.0, {}
274     # ===== #
275     # YOUR CODE HERE:
276     # Implement the backwards pass of the FC net and store the gradients
277     # in the grads dict, so that grads[k] is the gradient of self.params[k]
278     # Be sure your L2 regularization includes a 0.5 factor.
279     # ===== #
280
281     sm_loss, dout = softmax_loss(scores, y)

```

```

282     reg_loss = 0.0
283     for i in range(1, self.num_layers + 1):
284         reg_loss += 0.5 * self.reg * (np.sum(self.params['W{}'.format(i)]**2))
285     loss = sm_loss + reg_loss
286
287     for i in range(self.num_layers, 0, -1):
288         cur_w = 'W{}'.format(i)
289         cur_b = 'b{}'.format(i)
290         if (i == self.num_layers):
291             dout, grads[cur_w], grads[cur_b] = affine_backward(dout,
292 caches.pop())
293         else:
294             dout, grads[cur_w], grads[cur_b] = affine_relu_backward(dout,
295 caches.pop())
296             grads[cur_w] += self.reg * self.params[cur_w]
297
298     # ===== #
299     # END YOUR CODE HERE
300     # ===== #
301     return loss, grads

```