


1. a.) This minimization aims to make the difference between  $W^T W x$  and  $x$ , so the difference between  $W x$  and  $x$  will be minimized.

b.)




- c.) We account for these two paths by summing the upstream gradients, per the Law of total derivatives.

$$d.) \quad \frac{\partial \mathcal{L}}{\partial a} = a = \frac{\partial \mathcal{L}}{\partial b}$$

$$\frac{\partial \mathcal{L}}{\partial W^T} = \frac{\partial \mathcal{L}}{\partial b} \cdot c^T = a \cdot (W x)^T$$

~~$$\frac{\partial \mathcal{L}}{\partial c} = W \cdot \frac{\partial \mathcal{L}}{\partial b} = W^T \cdot a \quad \frac{\partial \mathcal{L}}{\partial c} = W \cdot \frac{\partial \mathcal{L}}{\partial b} = W \cdot a$$~~

$$\frac{\partial \mathcal{L}}{\partial W} = \frac{\partial \mathcal{L}}{\partial c} \cdot x^T = W a x^T$$



$$\nabla_W \mathcal{L} = W(W^T W x - x)x^T + (W^T W x - x)(W x)^T$$

# two\_layer\_nn

February 6, 2020

## 0.1 This is the 2-layer neural network workbook for ECE 239AS Assignment #3

Please follow the notebook linearly to implement a two layer neural network.

Please print out the workbook entirely when completed.

We thank Serena Yeung & Justin Johnson for permission to use code written for the CS 231n class ([cs231n.stanford.edu](http://cs231n.stanford.edu)). These are the functions in the cs231n folders and code in the jupyter notebook to preprocess and show the images. The classifiers used are based off of code prepared for CS 231n as well.

The goal of this workbook is to give you experience with training a two layer neural network.

```
In [1]: import random
import numpy as np
from cs231n.data_utils import load_CIFAR10
import matplotlib.pyplot as plt

%matplotlib inline
%load_ext autoreload
%autoreload 2

def rel_error(x, y):
    """ returns relative error """
    return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))
```

## 0.2 Toy example

Before loading CIFAR-10, there will be a toy example to test your implementation of the forward and backward pass

```
In [2]: from nndl.neural_net import TwoLayerNet
```

```
In [3]: # Create a small net and some toy data to check your implementations.
# Note that we set the random seed for repeatable experiments.

input_size = 4
hidden_size = 10
num_classes = 3
num_inputs = 5
```

```

def init_toy_model():
    np.random.seed(0)
    return TwoLayerNet(input_size, hidden_size, num_classes, std=1e-1)

def init_toy_data():
    np.random.seed(1)
    X = 10 * np.random.randn(num_inputs, input_size)
    y = np.array([0, 1, 2, 2, 1])
    return X, y

net = init_toy_model()
X, y = init_toy_data()

```

## 0.2.1 Compute forward pass scores

In [4]: *## Implement the forward pass of the neural network.*

```

# Note, there is a statement if y is None: return scores, which is why
# the following call will calculate the scores.
scores = net.loss(X)
print('Your scores:')
print(scores)
print()
print('correct scores:')
correct_scores = np.asarray([
    [-1.07260209,  0.05083871, -0.87253915],
    [-2.02778743, -0.10832494, -1.52641362],
    [-0.74225908,  0.15259725, -0.39578548],
    [-0.38172726,  0.10835902, -0.17328274],
    [-0.64417314, -0.18886813, -0.41106892]])
print(correct_scores)
print()

# The difference should be very small. We get < 1e-7
print('Difference between your scores and correct scores:')
print(np.sum(np.abs(scores - correct_scores)))

```

Your scores:

```

[[-1.07260209  0.05083871 -0.87253915]
 [-2.02778743 -0.10832494 -1.52641362]
 [-0.74225908  0.15259725 -0.39578548]
 [-0.38172726  0.10835902 -0.17328274]
 [-0.64417314 -0.18886813 -0.41106892]]

```

correct scores:

```

[[-1.07260209  0.05083871 -0.87253915]
 [-2.02778743 -0.10832494 -1.52641362]
 [-0.74225908  0.15259725 -0.39578548]]

```

```
[-0.38172726  0.10835902 -0.17328274]
[-0.64417314 -0.18886813 -0.41106892]]
```

Difference between your scores and correct scores:  
3.381231233889892e-08

## 0.2.2 Forward pass loss

```
In [5]: loss, _ = net.loss(X, y, reg=0.05)
        correct_loss = 1.071696123862817

        # should be very small, we get < 1e-12
        print("Loss:", loss)
        print('Difference between your loss and correct loss:')
        print(np.sum(np.abs(loss - correct_loss)))
```

Loss: 1.071696123862817  
Difference between your loss and correct loss:  
0.0

## 0.2.3 Backward pass

Implements the backwards pass of the neural network. Check your gradients with the gradient check utilities provided.

```
In [6]: from cs231n.gradient_check import eval_numerical_gradient

        # Use numeric gradient checking to check your implementation of the backward pass.
        # If your implementation is correct, the difference between the numeric and
        # analytic gradients should be less than 1e-8 for each of W1, W2, b1, and b2.

        loss, grads = net.loss(X, y, reg=0.05)

        # these should all be less than 1e-8 or so
        for param_name in grads:
            f = lambda W: net.loss(X, y, reg=0.05)[0]
            param_grad_num = eval_numerical_gradient(f, net.params[param_name], verbose=False)
            print('{} max relative error: {}'.format(param_name, rel_error(param_grad_num, grads[param_name])))
```

W2 max relative error: 2.9632227682005116e-10  
b2 max relative error: 1.2482651595953946e-09  
W1 max relative error: 1.2832908996874818e-09  
b1 max relative error: 3.1726798997101967e-09



### 0.2.4 Training the network


Implement `neural_net.train()` to train the network via stochastic gradient descent, much like the softmax and SVM.

```
In [7]: net = init_toy_model()
        stats = net.train(X, y, X, y,
                           learning_rate=1e-1, reg=5e-6,
                           num_iters=100, verbose=False)

        print('Final training loss: ', stats['loss_history'][-1])

        # plot the loss history
        plt.plot(stats['loss_history'])
        plt.xlabel('iteration')
        plt.ylabel('training loss')
        plt.title('Training Loss history')
        plt.show()
```

Final training loss: 0.014497864587765884



## 0.3 Classify CIFAR-10

Do classification on the CIFAR-10 dataset.

```

In [8]: from cs231n.data_utils import load_CIFAR10

def get_CIFAR10_data(num_training=49000, num_validation=1000, num_test=1000):
    """
    Load the CIFAR-10 dataset from disk and perform preprocessing to prepare
    it for the two-layer neural net classifier. These are the same steps as
    we used for the SVM, but condensed to a single function.
    """
    # Load the raw CIFAR-10 data
    cifar10_dir = './cifar-10-batches-py'
    X_train, y_train, X_test, y_test = load_CIFAR10(cifar10_dir)

    # Subsample the data
    mask = list(range(num_training, num_training + num_validation))
    X_val = X_train[mask]
    y_val = y_train[mask]
    mask = list(range(num_training))
    X_train = X_train[mask]
    y_train = y_train[mask]
    mask = list(range(num_test))
    X_test = X_test[mask]
    y_test = y_test[mask]

    # Normalize the data: subtract the mean image
    mean_image = np.mean(X_train, axis=0)
    X_train -= mean_image
    X_val -= mean_image
    X_test -= mean_image

    # Reshape data to rows
    X_train = X_train.reshape(num_training, -1)
    X_val = X_val.reshape(num_validation, -1)
    X_test = X_test.reshape(num_test, -1)

    return X_train, y_train, X_val, y_val, X_test, y_test

# Invoke the above function to get our data.
X_train, y_train, X_val, y_val, X_test, y_test = get_CIFAR10_data()
print('Train data shape: ', X_train.shape)
print('Train labels shape: ', y_train.shape)
print('Validation data shape: ', X_val.shape)
print('Validation labels shape: ', y_val.shape)
print('Test data shape: ', X_test.shape)
print('Test labels shape: ', y_test.shape)

```

```

Train data shape: (49000, 3072)
Train labels shape: (49000,)

```

```
Validation data shape: (1000, 3072)
Validation labels shape: (1000,)
Test data shape: (1000, 3072)
Test labels shape: (1000,)
```

### 0.3.1 Running SGD

If your implementation is correct, you should see a validation accuracy of around 28-29%.

```
In [9]: input_size = 32 * 32 * 3
        hidden_size = 50
        num_classes = 10
        net = TwoLayerNet(input_size, hidden_size, num_classes)

        # Train the network
        stats = net.train(X_train, y_train, X_val, y_val,
                           num_iters=1000, batch_size=200,
                           learning_rate=1e-4, learning_rate_decay=0.95,
                           reg=0.25, verbose=True)

        # Predict on the validation set
        val_acc = (net.predict(X_val) == y_val).mean()
        print('Validation accuracy: ', val_acc)

        # Save this net as the variable subopt_net for later comparison.
        subopt_net = net

iteration 0 / 1000: loss 2.302757518613176
iteration 100 / 1000: loss 2.302120159207236
iteration 200 / 1000: loss 2.2956136007408703
iteration 300 / 1000: loss 2.2518259043164135
iteration 400 / 1000: loss 2.188995235046776
iteration 500 / 1000: loss 2.1162527791897747
iteration 600 / 1000: loss 2.064670827698217
iteration 700 / 1000: loss 1.9901688623083942
iteration 800 / 1000: loss 2.002827640124685
iteration 900 / 1000: loss 1.9465176817856495
Validation accuracy: 0.283
```

## 0.4 Questions:


The training accuracy isn't great.

- (1) What are some of the reasons why this is the case? Take the following cell to do some analyses and then report your answers in the cell following the one below.
- (2) How should you fix the problems you identified in (1)?


```
In [10]: stats['train_acc_history']
```

```
Out[10]: [0.095, 0.15, 0.25, 0.25, 0.315]
```

```
In [11]: # ===== #  
# YOUR CODE HERE:  
# Do some debugging to gain some insight into why the optimization  
# isn't great.  
# ===== #  
  
# Plot the loss function and train / validation accuracies  
plt.plot(stats['train_acc_history'], label='Training Accuracy')  
plt.xlabel('Epochs')  
plt.ylabel('Accuracy')  
plt.plot(stats['val_acc_history'], label='Validation Accuracy')  
plt.legend()  
plt.show()  
  
plt.plot(stats['loss_history'])  
plt.show()  
  
# ===== #  
# END YOUR CODE HERE  
# ===== #
```







## 0.5 Answers:

- (1) Our learning rate seems to be zigzagging, which indicates a learning rate that is too large. However, since the training and validation accuracies are still closely aligned, we may not be training long enough.
- (2) We may want to increase our learning rate decay or decrease our learning rate so that our loss function converges faster. We may also want to train for more iterations, since it seems that we have not yet overfit our training data.

## 0.6 Optimize the neural network

Use the following part of the Jupyter notebook to optimize your hyperparameters on the validation set. Store your nets as `best_net`.

```
In [16]: best_net = None # store the best model into this

# ===== #
# YOUR CODE HERE:
# Optimize over your hyperparameters to arrive at the best neural
# network. You should be able to get over 50% validation accuracy.
# For this part of the notebook, we will give credit based on the
# accuracy you get. Your score on this question will be multiplied by:
# min(floor((X - 28%)) / %22, 1)
# where if you get 50% or higher validation accuracy, you get full
# points.
```

```

#
# Note, you need to use the same network structure (keep hidden_size = 50)!
# ===== #
input_size = 32 * 32 * 3
hidden_size = 50
num_classes = 10
best_net = TwoLayerNet(input_size, hidden_size, num_classes)

# Train the network
stats = best_net.train(X_train, y_train, X_val, y_val,
                       num_iters=4000, batch_size=200,
                       learning_rate=3e-3, learning_rate_decay=0.85,
                       reg=0.3, verbose=True)

# ===== #
# END YOUR CODE HERE
# ===== #
val_acc = (best_net.predict(X_val) == y_val).mean()
print('Validation accuracy: ', val_acc)

plt.plot(stats['train_acc_history'], label='Training Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.plot(stats['val_acc_history'], label='Validation Accuracy')
plt.legend()
plt.show()

plt.plot(stats['loss_history'])
plt.show()


```


```

iteration 0 / 4000: loss 2.3028326000608694
iteration 100 / 4000: loss 1.7255137664245441
iteration 200 / 4000: loss 1.5987255096449478
iteration 300 / 4000: loss 1.6251357161254112
iteration 400 / 4000: loss 1.6813738338603605
iteration 500 / 4000: loss 1.6005108953131277
iteration 600 / 4000: loss 1.5403566086755665
iteration 700 / 4000: loss 1.5567015543450125
iteration 800 / 4000: loss 1.5114857974651013
iteration 900 / 4000: loss 1.4129962223808301
iteration 1000 / 4000: loss 1.2735448943341665
iteration 1100 / 4000: loss 1.4279295575265616
iteration 1200 / 4000: loss 1.4239563180338874
iteration 1300 / 4000: loss 1.417309538366283
iteration 1400 / 4000: loss 1.3657055588859186
iteration 1500 / 4000: loss 1.3553863772909602
iteration 1600 / 4000: loss 1.4842079929712755
iteration 1700 / 4000: loss 1.3382454110453572

```

iteration 1800 / 4000: loss 1.4665457280699041  
iteration 1900 / 4000: loss 1.404499775614105  
iteration 2000 / 4000: loss 1.362044664325281  
iteration 2100 / 4000: loss 1.34841277411235  
iteration 2200 / 4000: loss 1.4236138060594488  
iteration 2300 / 4000: loss 1.4319267633413058  
iteration 2400 / 4000: loss 1.4015114433330438  
iteration 2500 / 4000: loss 1.2788842441558186  
iteration 2600 / 4000: loss 1.4617135866408193  
iteration 2700 / 4000: loss 1.33619618650787  
iteration 2800 / 4000: loss 1.352452528865135  
iteration 2900 / 4000: loss 1.3595451738246205  
iteration 3000 / 4000: loss 1.2661051686343472  
iteration 3100 / 4000: loss 1.2739861712687424  
iteration 3200 / 4000: loss 1.1324470048396282  
iteration 3300 / 4000: loss 1.3515679495238333  
iteration 3400 / 4000: loss 1.3456687150577051  
iteration 3500 / 4000: loss 1.296755025350978  
iteration 3600 / 4000: loss 1.3397932352835276  
iteration 3700 / 4000: loss 1.215692845056337  
iteration 3800 / 4000: loss 1.3483471620298124  
iteration 3900 / 4000: loss 1.1563534420779964  
Validation accuracy: 0.517







```
In [18]: from cs231n.vis_utils import visualize_grid
```

```
# Visualize the weights of the network
```

```
def show_net_weights(net):
    W1 = net.params['W1']
    W1 = W1.T.reshape(32, 32, 3, -1).transpose(3, 0, 1, 2)
    plt.imshow(visualize_grid(W1, padding=3).astype('uint8'))
    plt.gca().axis('off')
    plt.show()
```

```
show_net_weights(subopt_net)
```

```
show_net_weights(best_net)
```



## 0.7 Question:

- (1) What differences do you see in the weights between the suboptimal net and the best net you arrived at?

## 0.8 Answer:

- (1) The suboptimal net has much less complex weights than the best net I arrived at.

## 0.9 Evaluate on test set

```
In [19]: test_acc = (best_net.predict(X_test) == y_test).mean()  
         print('Test accuracy: ', test_acc)
```

Test accuracy: 0.528



```

1 import numpy as np
2 import matplotlib.pyplot as plt
3
4 """
5 This code was originally written for CS 231n at Stanford University
6 (cs231n.stanford.edu). It has been modified in various areas for use in the
7 ECE 239AS class at UCLA. This includes the descriptions of what code to
8 implement as well as some slight potential changes in variable names to be
9 consistent with class nomenclature. We thank Justin Johnson & Serena Yeung
10 for
11 permission to use this code. To see the original version, please visit
12 cs231n.stanford.edu.
13 """
14
15 class TwoLayerNet(object):
16     """
17     A two-layer fully-connected neural network. The net has an input dimension
18     of
19     N, a hidden layer dimension of H, and performs classification over C
20     classes.
21     We train the network with a softmax loss function and L2 regularization on the
22     weight matrices. The network uses a ReLU nonlinearity after the first fully
23     connected layer.
24
25     In other words, the network has the following architecture:
26
27     input - fully connected layer - ReLU - fully connected layer - softmax
28
29     The outputs of the second fully-connected layer are the scores for each
30     class.
31     """
32
33     def __init__(self, input_size, hidden_size, output_size, std=1e-4):
34         """
35         Initialize the model. Weights are initialized to small random values and
36         biases are initialized to zero. Weights and biases are stored in the
37         variable self.params, which is a dictionary with the following keys:
38
39         W1: First layer weights; has shape (H, D)
40         b1: First layer biases; has shape (H,)
41         W2: Second layer weights; has shape (C, H)
42         b2: Second layer biases; has shape (C,)
43
44         Inputs:
45         - input_size: The dimension D of the input data.
46         - hidden_size: The number of neurons H in the hidden layer.
47         - output_size: The number of classes C.
48         """
49         self.params = {}
50         self.params['W1'] = std * np.random.randn(hidden_size, input_size)
51         self.params['b1'] = np.zeros(hidden_size)
52         self.params['W2'] = std * np.random.randn(output_size, hidden_size)
53         self.params['b2'] = np.zeros(output_size)
54
55     def loss(self, X, y=None, reg=0.0):
56         """
57         Compute the loss and gradients for a two layer fully connected neural
58         network.

```

```

56
57 Inputs:
58 - X: Input data of shape (N, D). Each X[i] is a training sample.
59 - y: Vector of training labels. y[i] is the label for X[i], and each y[i]
is
60 an integer in the range  $0 \leq y[i] < C$ . This parameter is optional; if
it
61 is not passed then we only return scores, and if it is passed then we
62 instead return the loss and gradients.
63 - reg: Regularization strength.
64
65 Returns:
66 If y is None, return a matrix scores of shape (N, C) where scores[i, c]
is
67 the score for class c on input X[i].
68
69 If y is not None, instead return a tuple of:
70 - loss: Loss (data loss and regularization loss) for this batch of
training
71 samples.
72 - grads: Dictionary mapping parameter names to gradients of those
parameters
73 with respect to the loss function; has the same keys as self.params.
74 """
75 # Unpack variables from the params dictionary
76 W1, b1 = self.params['W1'], self.params['b1']
77 W2, b2 = self.params['W2'], self.params['b2']
78 N, D = X.shape
79
80 # Compute the forward pass
81 scores = None
82
83 # ===== #
84 # YOUR CODE HERE:
85 # Calculate the output scores of the neural network. The result
86 # should be (N, C). As stated in the description for this class,
87 # there should not be a ReLU layer after the second FC layer.
88 # The output of the second FC layer is the output scores. Do not
89 # use a for loop in your implementation.
90 # ===== #
91
92 h1 = np.maximum(X.dot(W1.T) + b1, 0)
93 h2 = h1.dot(W2.T) + b2
94 scores = h2
95
96 # ===== #
97 # END YOUR CODE HERE
98 # ===== #
99
100 # If the targets are not given then jump out, we're done
101 if y is None:
102     return scores
103
104 # Compute the loss
105 loss = None
106
107 # ===== #
108 # YOUR CODE HERE:
109 # Calculate the loss of the neural network. This includes the
110 # softmax loss and the L2 regularization for W1 and W2. Store the

```

```

111 # total loss in the variable loss. Multiply the regularization
112 #   loss by 0.5 (in addition to the factor reg).
113 # ===== #
114
115     # scores is num_examples by num_classes
116     ea = np.exp(scores - np.max(scores))
117     sums = np.sum(ea, axis=1)
118     softmax = ea / sums[:, np.newaxis]
119     loss = np.mean(-np.log(softmax[np.arange(N), y]))
120     loss += 0.5 * reg * (np.sum(W1**2) + np.sum(W2**2))
121     # ===== #
122     # END YOUR CODE HERE
123     # ===== #
124
125     grads = {}
126
127     # ===== #
128     # YOUR CODE HERE:
129 # Implement the backward pass. Compute the derivatives of the
130 # weights and the biases. Store the results in the grads
131 # dictionary. e.g., grads['W1'] should store the gradient for
132 # W1, and be of the same size as W1.
133 # ===== #
134
135     grad_softmax = softmax
136     grad_softmax[np.arange(N), y] -= 1
137     grad_softmax /= N
138     grads['W2'] = grad_softmax.T.dot(h1) + reg * W2    # (C x N) * (N x H) =
(C x H)
139     grads['b2'] = np.sum(grad_softmax, axis=0)        # (C, )
140     # (H x C) * (C x N) = (H x N)
141     # (H x N) * (N x D) = (H x D)
142     ind = h1
143     ind[ind > 0] = 1
144     ind[ind <= 0] = 0
145     grads['W1'] = (W2.T.dot(grad_softmax.T) * ind.T).dot(X) + reg * W1
146     grads['b1'] = np.sum((W2.T.dot(grad_softmax.T) * ind.T).T, axis = 0)
147
148     # ===== #
149     # END YOUR CODE HERE
150     # ===== #
151
152     return loss, grads
153
154 def train(self, X, y, X_val, y_val,
155           learning_rate=1e-3, learning_rate_decay=0.95,
156           reg=1e-5, num_iters=100,
157           batch_size=200, verbose=False):
158     """
159     Train this neural network using stochastic gradient descent.
160
161     Inputs:
162     - X: A numpy array of shape (N, D) giving training data.
163     - y: A numpy array of shape (N,) giving training labels; y[i] = c means
that
164         X[i] has label c, where 0 <= c < C.
165     - X_val: A numpy array of shape (N_val, D) giving validation data.
166     - y_val: A numpy array of shape (N_val,) giving validation labels.
167     - learning_rate: Scalar giving learning rate for optimization.

```

```

168     - learning_rate_decay: Scalar giving factor used to decay the learning
rate
169         after each epoch.
170     - reg: Scalar giving regularization strength.
171     - num_iters: Number of steps to take when optimizing.
172     - batch_size: Number of training examples to use per step.
173     - verbose: boolean; if true print progress during optimization.
174     """
175     num_train = X.shape[0]
176     iterations_per_epoch = max(num_train / batch_size, 1)
177
178     # Use SGD to optimize the parameters in self.model
179     loss_history = []
180     train_acc_history = []
181     val_acc_history = []
182
183     for it in np.arange(num_iters):
184         X_batch = None
185         y_batch = None
186
187         # ===== #
188         # YOUR CODE HERE:
189     # Create a minibatch by sampling batch_size samples randomly.
190     # ===== #
191         indices = np.random.choice(num_train, batch_size)
192         X_batch = X[indices]
193         y_batch = y[indices]
194
195         # ===== #
196         # END YOUR CODE HERE
197         # ===== #
198
199         # Compute loss and gradients using the current minibatch
200         loss, grads = self.loss(X_batch, y=y_batch, reg=reg)
201         loss_history.append(loss)
202
203         # ===== #
204         # YOUR CODE HERE:
205     # Perform a gradient descent step using the minibatch to update
206     # all parameters (i.e., W1, W2, b1, and b2).
207     # ===== #
208
209         self.params['W1'] -= learning_rate * grads['W1']
210         self.params['b1'] -= learning_rate * grads['b1']
211         self.params['W2'] -= learning_rate * grads['W2']
212         self.params['b2'] -= learning_rate * grads['b2']
213
214         # ===== #
215         # END YOUR CODE HERE
216         # ===== #
217
218         if verbose and it % 100 == 0:
219             print('iteration {} / {}: loss {}'.format(it, num_iters, loss))
220
221         # Every epoch, check train and val accuracy and decay learning rate.
222         if it % iterations_per_epoch == 0:
223             # Check accuracy
224             train_acc = (self.predict(X_batch) == y_batch).mean()
225             val_acc = (self.predict(X_val) == y_val).mean()
226             train_acc_history.append(train_acc)

```

```

227         val_acc_history.append(val_acc)
228
229         # Decay learning rate
230         learning_rate *= learning_rate_decay
231
232     return {
233         'loss_history': loss_history,
234         'train_acc_history': train_acc_history,
235         'val_acc_history': val_acc_history,
236     }
237
238     def predict(self, X):
239         """
240         Use the trained weights of this two-layer network to predict labels for
241         data points. For each data point we predict scores for each of the C
242         classes, and assign each data point to the class with the highest score.
243
244         Inputs:
245         - X: A numpy array of shape (N, D) giving N D-dimensional data points to
246             classify.
247
248         Returns:
249         - y_pred: A numpy array of shape (N,) giving predicted labels for each of
250             the elements of X. For all i, y_pred[i] = c means that X[i] is
251             predicted
252             to have class c, where 0 <= c < C.
253         """
254         y_pred = None
255
256         # ===== #
257         # YOUR CODE HERE:
258         # Predict the class given the input data.
259         # ===== #
260
261         scores = np.maximum(X.dot(self.params['W1']).T + self.params['b1'],
262                             0).dot(self.params['W2']).T + self.params['b2']
263         y_pred = np.argmax(scores, axis=1)
264
265         # ===== #
266         # END YOUR CODE HERE
267         # ===== #
268
269         return y_pred

```

# FC\_nets

February 6, 2020

## 1 Fully connected networks

In the previous notebook, you implemented a simple two-layer neural network class. However, this class is not modular. If you wanted to change the number of layers, you would need to write a new loss and gradient function. If you wanted to optimize the network with different optimizers, you'd need to write new training functions. If you wanted to incorporate regularizations, you'd have to modify the loss and gradient function.

Instead of having to modify functions each time, for the rest of the class, we'll work in a more modular framework where we define forward and backward layers that calculate losses and gradients respectively. Since the forward and backward layers share intermediate values that are useful for calculating both the loss and the gradient, we'll also have these function return "caches" which store useful intermediate values.

The goal is that through this modular design, we can build different sized neural networks for various applications.

In this HW #3, we'll define the basic architecture, and in HW #4, we'll build on this framework to implement different optimizers and regularizations (like BatchNorm and Dropout).

CS231n has built a solid API for building these modular frameworks and training them, and we will use their very well implemented framework as opposed to "reinventing the wheel." This includes using their Solver, various utility functions, and their layer structure. This also includes `nndl.fc_net`, `nndl.layers`, and `nndl.layer_utils`. As in prior assignments, we thank Serena Yeung & Justin Johnson for permission to use code written for the CS 231n class ([cs231n.stanford.edu](http://cs231n.stanford.edu)).

### 1.1 Modular layers

This notebook will build modular layers in the following manner. First, there will be a forward pass for a given layer with inputs (`x`) and return the output of that layer (`out`) as well as cached variables (`cache`) that will be used to calculate the gradient in the backward pass.

```
def layer_forward(x, w):
    """ Receive inputs x and weights w """
    # Do some computations ...
    z = # ... some intermediate value
    # Do some more computations ...
    out = # the output

    cache = (x, w, z, out) # Values we need to compute gradients

    return out, cache
```



The backward pass will receive upstream derivatives and the cache object, and will return gradients with respect to the inputs and weights, like this:

```
def layer_backward(dout, cache):
    """
    Receive derivative of loss with respect to outputs and cache,
    and compute derivative with respect to inputs.
    """
    # Unpack cache values
    x, w, z, out = cache

    # Use values in cache to compute derivatives
    dx = # Derivative of loss with respect to x
    dw = # Derivative of loss with respect to w

    return dx, dw

In [1]: ## Import and setups

import time
import numpy as np
import matplotlib.pyplot as plt
from nndl.fc_net import *
from cs231n.data_utils import get_CIFAR10_data
from cs231n.gradient_check import eval_numerical_gradient, eval_numerical_gradient_array
from cs231n.solver import Solver

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2

def rel_error(x, y):
    """ returns relative error """
    return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))

In [2]: # Load the (preprocessed) CIFAR10 data.

data = get_CIFAR10_data()
for k in data.keys():
    print('{}: {}'.format(k, data[k].shape))

X_train: (49000, 3, 32, 32)
y_train: (49000,)
```

```
X_val: (1000, 3, 32, 32)
y_val: (1000,)
X_test: (1000, 3, 32, 32)
y_test: (1000,)
```

## 1.2 Linear layers

In this section, we'll implement the forward and backward pass for the linear layers.

The linear layer forward pass is the function `affine_forward` in `nn1/layers.py` and the backward pass is `affine_backward`.

After you have implemented these, test your implementation by running the cell below.

### 1.2.1 Affine layer forward pass

Implement `affine_forward` and then test your code by running the following cell.

```
In [3]: # Test the affine_forward function
```

```
num_inputs = 2
input_shape = (4, 5, 6)
output_dim = 3

input_size = num_inputs * np.prod(input_shape)
weight_size = output_dim * np.prod(input_shape)

x = np.linspace(-0.1, 0.5, num=input_size).reshape(num_inputs, *input_shape)
w = np.linspace(-0.2, 0.3, num=weight_size).reshape(np.prod(input_shape), output_dim)
b = np.linspace(-0.3, 0.1, num=output_dim)

out, _ = affine_forward(x, w, b)
correct_out = np.array([[ 1.49834967,  1.70660132,  1.91485297],
                        [ 3.25553199,  3.5141327,  3.77273342]])

# Compare your output with ours. The error should be around 1e-9.
print('Testing affine_forward function:')
print('difference: {}'.format(rel_error(out, correct_out)))
```

```
Testing affine_forward function:
difference: 9.769849468192957e-10
```

### 1.2.2 Affine layer backward pass

Implement `affine_backward` and then test your code by running the following cell.

```
In [4]: # Test the affine_backward function
```

```
x = np.random.randn(10, 2, 3)
```

```

w = np.random.randn(6, 5)
b = np.random.randn(5)
dout = np.random.randn(10, 5)

dx_num = eval_numerical_gradient_array(lambda x: affine_forward(x, w, b)[0], x, dout)
dw_num = eval_numerical_gradient_array(lambda w: affine_forward(x, w, b)[0], w, dout)
db_num = eval_numerical_gradient_array(lambda b: affine_forward(x, w, b)[0], b, dout)

_, cache = affine_forward(x, w, b)
dx, dw, db = affine_backward(dout, cache)

# The error should be around 1e-10
print('Testing affine_backward function:')
print('dx error: {}'.format(rel_error(dx_num, dx)))
print('dw error: {}'.format(rel_error(dw_num, dw)))
print('db error: {}'.format(rel_error(db_num, db)))

```

```

Testing affine_backward function:
dx error: 3.771222131555466e-10
dw error: 3.4224790236683927e-09
db error: 9.60656250860874e-12

```

## 1.3 Activation layers

In this section you'll implement the ReLU activation.

### 1.3.1 ReLU forward pass

Implement the `relu_forward` function in `nndl/layers.py` and then test your code by running the following cell.

```

In [5]: # Test the relu_forward function

x = np.linspace(-0.5, 0.5, num=12).reshape(3, 4)

out, _ = relu_forward(x)
correct_out = np.array([[ 0.,          0.,          0.,          0.,          ],
                        [ 0.,          0.,          0.04545455, 0.13636364, ],
                        [ 0.22727273, 0.31818182, 0.40909091, 0.5,          ]])

# Compare your output with ours. The error should be around 1e-8
print('Testing relu_forward function:')
print('difference: {}'.format(rel_error(out, correct_out)))

```

```

Testing relu_forward function:
difference: 4.999999798022158e-08

```

### 1.3.2 ReLU backward pass

Implement the `relu_backward` function in `nndl/layers.py` and then test your code by running the following cell.

```
In [6]: x = np.random.randn(10, 10)
        dout = np.random.randn(*x.shape)

        dx_num = eval_numerical_gradient_array(lambda x: relu_forward(x)[0], x, dout)

        _, cache = relu_forward(x)
        dx = relu_backward(dout, cache)

        # The error should be around 1e-12
        print('Testing relu_backward function:')
        print('dx error: {}'.format(rel_error(dx_num, dx)))
```

```
Testing relu_backward function:
dx error: 3.2755894889487494e-12
```

## 1.4 Combining the affine and ReLU layers

Often times, an affine layer will be followed by a ReLU layer. So let's make one that puts them together. Layers that are combined are stored in `nndl/layer_utils.py`.

### 1.4.1 Affine-ReLU layers

We've implemented `affine_relu_forward()` and `affine_relu_backward` in `nndl/layer_utils.py`. Take a look at them to make sure you understand what's going on. Then run the following cell to ensure its implemented correctly.

```
In [7]: from nndl.layer_utils import affine_relu_forward, affine_relu_backward

        x = np.random.randn(2, 3, 4)
        w = np.random.randn(12, 10)
        b = np.random.randn(10)
        dout = np.random.randn(2, 10)

        out, cache = affine_relu_forward(x, w, b)
        dx, dw, db = affine_relu_backward(dout, cache)

        dx_num = eval_numerical_gradient_array(lambda x: affine_relu_forward(x, w, b)[0], x, dout)
        dw_num = eval_numerical_gradient_array(lambda w: affine_relu_forward(x, w, b)[0], w, dout)
        db_num = eval_numerical_gradient_array(lambda b: affine_relu_forward(x, w, b)[0], b, dout)

        print('Testing affine_relu_forward and affine_relu_backward:')
        print('dx error: {}'.format(rel_error(dx_num, dx)))
        print('dw error: {}'.format(rel_error(dw_num, dw)))
        print('db error: {}'.format(rel_error(db_num, db)))
```

```
Testing affine_relu_forward and affine_relu_backward:
dx error: 1.1103532464208299e-10
dw error: 1.7216726870718686e-10
db error: 7.826682727335577e-12
```

## 1.5 Softmax and SVM losses

You've already implemented these, so we have written these in `layers.py`. The following code will ensure they are working correctly.

```
In [8]: num_classes, num_inputs = 10, 50
        x = 0.001 * np.random.randn(num_inputs, num_classes)
        y = np.random.randint(num_classes, size=num_inputs)

        dx_num = eval_numerical_gradient(lambda x: svm_loss(x, y)[0], x, verbose=False)
        loss, dx = svm_loss(x, y)

        # Test svm_loss function. Loss should be around 9 and dx error should be 1e-9
        print('Testing svm_loss:')
        print('loss: {}'.format(loss))
        print('dx error: {}'.format(rel_error(dx_num, dx)))

        dx_num = eval_numerical_gradient(lambda x: softmax_loss(x, y)[0], x, verbose=False)
        loss, dx = softmax_loss(x, y)

        # Test softmax_loss function. Loss should be 2.3 and dx error should be 1e-8
        print('\nTesting softmax_loss:')
        print('loss: {}'.format(loss))
        print('dx error: {}'.format(rel_error(dx_num, dx)))
```

```
Testing svm_loss:
loss: 8.999851589061882
dx error: 1.4021566006651672e-09
```

```
Testing softmax_loss:
loss: 2.30257068534768
dx error: 7.987041003717542e-09
```

## 1.6 Implementation of a two-layer NN

In `nndl/fc_net.py`, implement the class `TwoLayerNet` which uses the layers you made here. When you have finished, the following cell will test your implementation.

```
In [9]: N, D, H, C = 3, 5, 50, 7
        X = np.random.randn(N, D)
        y = np.random.randint(C, size=N)
```

```

std = 1e-2
model = TwoLayerNet(input_dim=D, hidden_dims=H, num_classes=C, weight_scale=std)

print('Testing initialization ... ')
W1_std = abs(model.params['W1'].std() - std)
b1 = model.params['b1']
W2_std = abs(model.params['W2'].std() - std)
b2 = model.params['b2']
assert W1_std < std / 10, 'First layer weights do not seem right'
assert np.all(b1 == 0), 'First layer biases do not seem right'
assert W2_std < std / 10, 'Second layer weights do not seem right'
assert np.all(b2 == 0), 'Second layer biases do not seem right'

print('Testing test-time forward pass ... ')
model.params['W1'] = np.linspace(-0.7, 0.3, num=D*H).reshape(D, H)
model.params['b1'] = np.linspace(-0.1, 0.9, num=H)
model.params['W2'] = np.linspace(-0.3, 0.4, num=H*C).reshape(H, C)
model.params['b2'] = np.linspace(-0.9, 0.1, num=C)
X = np.linspace(-5.5, 4.5, num=N*D).reshape(D, N).T
scores = model.loss(X)
correct_scores = np.asarray(
    [[11.53165108, 12.2917344, 13.05181771, 13.81190102, 14.57198434, 15.33206765,
      12.05769098, 12.74614105, 13.43459113, 14.1230412, 14.81149128, 15.49994135,
      12.58373087, 13.20054771, 13.81736455, 14.43418138, 15.05099822, 15.66781506,
      12.25768488, 12.94613495, 13.63458502, 14.3230351, 15.01148517, 15.70003524,
      12.37767488, 13.06612495, 13.75457502, 14.4430251, 15.13147517, 15.82002524,
      12.50000000, 13.18750000, 13.87500000, 14.56250000, 15.25000000, 15.93750000,
      12.62500000, 13.31250000, 14.00000000, 14.68750000, 15.37500000, 16.06250000,
      12.75000000, 13.43750000, 14.12500000, 14.81250000, 15.50000000, 16.18750000,
      12.87500000, 13.56250000, 14.25000000, 14.93750000, 15.62500000, 16.31250000,
      13.00000000, 13.68750000, 14.37500000, 15.06250000, 15.75000000, 16.43750000,
      13.12500000, 13.81250000, 14.50000000, 15.18750000, 15.87500000, 16.56250000,
      13.25000000, 13.93750000, 14.62500000, 15.31250000, 16.00000000, 16.68750000,
      13.37500000, 14.06250000, 14.75000000, 15.43750000, 16.12500000, 16.81250000,
      13.50000000, 14.18750000, 14.87500000, 15.56250000, 16.25000000, 16.93750000,
      13.62500000, 14.31250000, 15.00000000, 15.68750000, 16.37500000, 17.06250000,
      13.75000000, 14.43750000, 15.12500000, 15.81250000, 16.50000000, 17.18750000,
      13.87500000, 14.56250000, 15.25000000, 15.93750000, 16.62500000, 17.31250000,
      14.00000000, 14.68750000, 15.37500000, 16.06250000, 16.75000000, 17.43750000,
      14.12500000, 14.81250000, 15.50000000, 16.18750000, 16.87500000, 17.56250000,
      14.25000000, 14.93750000, 15.62500000, 16.31250000, 17.00000000, 17.68750000,
      14.37500000, 15.06250000, 15.75000000, 16.43750000, 17.12500000, 17.81250000,
      14.50000000, 15.18750000, 15.87500000, 16.56250000, 17.25000000, 17.93750000,
      14.62500000, 15.31250000, 16.00000000, 16.68750000, 17.37500000, 18.06250000,
      14.75000000, 15.43750000, 16.12500000, 16.81250000, 17.50000000, 18.18750000,
      14.87500000, 15.56250000, 16.25000000, 16.93750000, 17.62500000, 18.31250000,
      15.00000000, 15.68750000, 16.37500000, 17.06250000, 17.75000000, 18.43750000,
      15.12500000, 15.81250000, 16.50000000, 17.18750000, 17.87500000, 18.56250000,
      15.25000000, 15.93750000, 16.62500000, 17.31250000, 18.00000000, 18.68750000,
      15.37500000, 16.06250000, 16.75000000, 17.43750000, 18.12500000, 18.81250000,
      15.50000000, 16.18750000, 16.87500000, 17.56250000, 18.25000000, 18.93750000,
      15.62500000, 16.31250000, 17.00000000, 17.68750000, 18.37500000, 19.06250000,
      15.75000000, 16.43750000, 17.12500000, 17.81250000, 18.50000000, 19.18750000,
      15.87500000, 16.56250000, 17.25000000, 17.93750000, 18.62500000, 19.31250000,
      16.00000000, 16.68750000, 17.37500000, 18.06250000, 18.75000000, 19.43750000,
      16.12500000, 16.81250000, 17.50000000, 18.18750000, 18.87500000, 19.56250000,
      16.25000000, 16.93750000, 17.62500000, 18.31250000, 19.00000000, 19.68750000,
      16.37500000, 17.06250000, 17.75000000, 18.43750000, 19.12500000, 19.81250000,
      16.50000000, 17.18750000, 17.87500000, 18.56250000, 19.25000000, 19.93750000,
      16.62500000, 17.31250000, 18.00000000, 18.68750000, 19.37500000, 20.06250000,
      16.75000000, 17.43750000, 18.12500000, 18.81250000, 19.50000000, 20.18750000,
      16.87500000, 17.56250000, 18.25000000, 18.93750000, 19.62500000, 20.31250000,
      17.00000000, 17.68750000, 18.37500000, 19.06250000, 19.75000000, 20.43750000,
      17.12500000, 17.81250000, 18.50000000, 19.18750000, 19.87500000, 20.56250000,
      17.25000000, 17.93750000, 18.62500000, 19.31250000, 20.00000000, 20.68750000,
      17.37500000, 18.06250000, 18.75000000, 19.43750000, 20.12500000, 20.81250000,
      17.50000000, 18.18750000, 18.87500000, 19.56250000, 20.25000000, 20.93750000,
      17.62500000, 18.31250000, 19.00000000, 19.68750000, 20.37500000, 21.06250000,
      17.75000000, 18.43750000, 19.12500000, 19.81250000, 20.50000000, 21.18750000,
      17.87500000, 18.56250000, 19.25000000, 19.93750000, 20.62500000, 21.31250000,
      18.00000000, 18.68750000, 19.37500000, 20.06250000, 20.75000000, 21.43750000,
      18.12500000, 18.81250000, 19.50000000, 20.18750000, 20.87500000, 21.56250000,
      18.25000000, 18.93750000, 19.62500000, 20.31250000, 21.00000000, 21.68750000,
      18.37500000, 19.06250000, 19.75000000, 20.43750000, 21.12500000, 21.81250000,
      18.50000000, 19.18750000, 19.87500000, 20.56250000, 21.25000000, 21.93750000,
      18.62500000, 19.31250000, 19.93750000, 20.68750000, 21.37500000, 22.06250000,
      18.75000000, 19.43750000, 20.00000000, 20.81250000, 21.50000000, 22.18750000,
      18.87500000, 19.56250000, 20.12500000, 20.93750000, 21.62500000, 22.31250000,
      19.00000000, 19.68750000, 20.25000000, 21.06250000, 21.75000000, 22.43750000,
      19.12500000, 19.81250000, 20.37500000, 21.18750000, 21.87500000, 22.56250000,
      19.25000000, 19.93750000, 20.50000000, 21.31250000, 22.00000000, 22.68750000,
      19.37500000, 20.06250000, 20.62500000, 21.43750000, 22.12500000, 22.81250000,
      19.50000000, 20.18750000, 20.75000000, 21.56250000, 22.25000000, 22.93750000,
      19.62500000, 20.31250000, 20.87500000, 21.68750000, 22.37500000, 23.06250000,
      19.75000000, 20.43750000, 21.00000000, 21.81250000, 22.50000000, 23.18750000,
      19.87500000, 20.56250000, 21.12500000, 21.93750000, 22.62500000, 23.31250000,
      19.93750000, 20.62500000, 21.18750000, 22.00000000, 22.68750000, 23.37500000,
      20.00000000, 20.68750000, 21.25000000, 22.06250000, 22.75000000, 23.43750000,
      20.06250000, 20.75000000, 21.31250000, 22.12500000, 22.81250000, 23.50000000,
      20.12500000, 20.81250000, 21.37500000, 22.18750000, 22.87500000, 23.56250000,
      20.18750000, 20.87500000, 21.43750000, 22.25000000, 22.93750000, 23.62500000,
      20.25000000, 20.93750000, 21.50000000, 22.31250000, 23.00000000, 23.68750000,
      20.31250000, 21.00000000, 21.56250000, 22.37500000, 23.06250000, 23.75000000,
      20.37500000, 21.06250000, 21.62500000, 22.43750000, 23.12500000, 23.81250000,
      20.43750000, 21.12500000, 21.68750000, 22.50000000, 23.18750000, 23.87500000,
      20.50000000, 21.18750000, 21.75000000, 22.56250000, 23.25000000, 23.93750000,
      20.56250000, 21.25000000, 21.81250000, 22.62500000, 23.31250000, 24.00000000,
      20.62500000, 21.31250000, 21.87500000, 22.68750000, 23.37500000, 24.06250000,
      20.68750000, 21.37500000, 21.93750000, 22.75000000, 23.43750000, 24.12500000,
      20.75000000, 21.43750000, 22.00000000, 22.81250000, 23.50000000, 24.18750000,
      20.81250000, 21.50000000, 22.06250000, 22.87500000, 23.56250000, 24.25000000,
      20.87500000, 21.56250000, 22.12500000, 22.93750000, 23.62500000, 24.31250000,
      20.93750000, 21.62500000, 22.18750000, 23.00000000, 23.68750000, 24.37500000,
      21.00000000, 21.68750000, 22.25000000, 23.06250000, 23.75000000, 24.43750000,
      21.06250000, 21.75000000, 22.31250000, 23.12500000, 23.81250000, 24.50000000,
      21.12500000, 21.81250000, 22.37500000, 23.18750000, 23.87500000, 24.56250000,
      21.18750000, 21.87500000, 22.43750000, 23.25000000, 23.93750000, 24.62500000,
      21.25000000, 21.93750000, 22.50000000, 23.31250000, 24.00000000, 24.68750000,
      21.31250000, 22.00000000, 22.56250000, 23.37500000, 24.06250000, 24.75000000,
      21.37500000, 22.06250000, 22.62500000, 23.43750000, 24.12500000, 24.81250000,
      21.43750000, 22.12500000, 22.68750000, 23.50000000, 24.18750000, 24.87500000,
      21.50000000, 22.18750000, 22.75000000, 23.56250000, 24.25000000, 24.93750000,
      21.56250000, 22.25000000, 22.81250000, 23.62500000, 24.31250000, 25.00000000,
      21.62500000, 22.31250000, 22.87500000, 23.68750000, 24.37500000, 25.06250000,
      21.68750000, 22.37500000, 22.93750000, 23.75000000, 24.43750000, 25.12500000,
      21.75000000, 22.43750000, 23.00000000, 23.81250000, 24.50000000, 25.18750000,
      21.81250000, 22.50000000, 23.06250000, 23.87500000, 24.56250000, 25.25000000,
      21.87500000, 22.56250000, 23.12500000, 23.93750000, 24.62500000, 25.31250000,
      21.93750000, 22.62500000, 23.18750000, 24.00000000, 24.68750000, 25.37500000,
      22.00000000, 22.68750000, 23.25000000, 24.06250000, 24.75000000, 25.43750000,
      22.06250000, 22.75000000, 23.31250000, 24.12500000, 24.81250000, 25.50000000,
      22.12500000, 22.81250000, 23.37500000, 24.18750000, 24.87500000, 25.56250000,
      22.18750000, 22.87500000, 23.43750000, 24.25000000, 24.93750000, 25.62500000,
      22.25000000, 22.93750000, 23.50000000, 24.31250000, 25.00000000, 25.68750000,
      22.31250000, 23.00000000, 23.56250000, 24.37500000, 25.06250000, 25.75000000,
      22.37500000, 23.06250000, 23.62500000, 24.43750000, 25.12500000, 25.81250000,
      22.43750000, 23.12500000, 23.68750000, 24.50000000, 25.18750000, 25.87500000,
      22.50000000, 23.18750000, 23.75000000, 24.56250000, 25.25000000, 25.93750000,
      22.56250000, 23.25000000, 23.81250000, 24.62500000, 25.31250000, 26.00000000,
      22.62500000, 23.31250000, 23.87500000, 24.68750000, 25.37500000, 26.06250000,
      22.68750000, 23.37500000, 23.93750000, 24.75000000, 25.43750000, 26.12500000,
      22.75000000, 23.43750000, 24.00000000, 24.81250000, 25.50000000, 26.18750000,
      22.81250000, 23.50000000, 24.06250000, 24.87500000, 25.56250000, 26.25000000,
      22.87500000, 23.56250000, 24.12500000, 24.93750000, 25.62500000, 26.31250000,
      22.93750000, 23.62500000, 24.18750000, 25.00000000, 25.68750000, 26.37500000,
      23.00000000, 23.68750000, 24.25000000, 25.06250000, 25.75000000, 26.43750000,
      23.06250000, 23.75000000, 24.31250000, 25.12500000, 25.81250000, 26.50000000,
      23.12500000, 23.81250000, 24.37500000, 25.18750000, 25.87500000, 26.56250000,
      23.18750000, 23.87500000, 24.43750000, 25.25000000, 25.93750000, 26.62500000,
      23.25000000, 23.93750000, 24.50000000, 25.31250000, 26.00000000, 26.68750000,
      23.31250000, 24.00000000, 24.56250000, 25.37500000, 26.06250000, 26.75000000,
      23.37500000, 24.06250000, 24.62500000, 25.43750000, 26.12500000, 26.81250000,
      23.43750000, 24.12500000, 24.68750000, 25.50000000, 26.18750000, 26.87500000,
      23.50000000, 24.18750000, 24.75000000, 25.56250000, 26.25000000, 26.93750000,
      23.56250000, 24.25000000, 24.81250000, 25.62500000, 26.31250000, 27.00000000,
      23.62500000, 24.31250000, 24.87500000, 25.68750000, 26.37500000, 27.06250000,
      23.68750000, 24.37500000, 24.93750000, 25.75000000, 26.43750000, 27.12500000,
      23.75000000, 24.43750000, 25.00000000, 25.81250000, 26.50000000, 27.18750000,
      23.81250000, 24.50000000, 25.06250000, 25.87500000, 26.56250000, 27.25000000,
      23.87500000, 24.56250000, 25.12500000, 25.93750000, 26.62500000, 27.31250000,
      23.93750000, 24.62500000, 25.18750000, 26.00000000, 26.68750000, 27.37500000,
      24.00000000, 24.68750000, 25.25000000, 26.06250000, 26.75000000, 27.43750000,
      24.06250000, 24.75000000, 25.31250000, 26.12500000, 26.81250000, 27.50000000,
      24.12500000, 24.81250000, 25.37500000, 26.18750000, 26.87500000, 27.56250000,
      24.18750000, 24.87500000, 25.43750000, 26.25000000, 26.93750000, 27.62500000,
      24.25000000, 24.93750000, 25.50000000, 26.31250000, 27.00000000, 27.68750000,
      24.31250000, 25.00000000, 25.56250000, 26.37500000, 27.06250000, 27.75000000,
      24.37500000, 25.06250000, 25.62500000, 26.43750000, 27.12500000, 27.81250000,
      24.43750000, 25.12500000, 25.68750000, 26.50000000, 27.18750000, 27.87500000,
      24.50000000, 25.18750000, 25.75000000, 26.56250000, 27.25000000, 27.93750000,
      24.56250000, 25.25000000, 25.81250000, 26.62500000, 27.31250000, 28.00000000,
      24.62500000, 25.31250000, 25.87500000, 26.68750000, 27.37500000, 28.06250000,
      24.68750000, 25.37500000, 25.93750000, 26.75000000, 27.43750000, 28.12500000,
      24.75000000, 25.43750000, 26.00000000, 26.81250000, 27.50000000, 28.18750000,
      24.81250000, 25.50000000, 26.06250000, 26.87500000, 27.56250000, 28.25000000,
      24.87500000, 25.56250000, 26.12500000, 26.93750000, 27.62500000, 28.31250000,
      24.93750000, 25.62500000, 26.18750000, 27.00000000, 27.68750000, 28.37500000,
      25.00000000, 25.68750000, 26.25000000, 27.06250000, 27.75000000, 28.43750000,
      25.06250000, 25.75000000, 26.31250000, 27.12500000, 27.81250000, 28.50000000,
      25.12500000, 25.81250000, 26.37500000, 27.18750000, 27.87500000, 28.56250000,
      25.18750000, 25.87500000, 26.43750000, 27.25000000, 27.93750000, 28.62500000,
      25.25000000, 25.93750000
```



```

Testing initialization ...
Testing test-time forward pass ...
Testing training loss (no regularization)
Running numeric gradient check with reg = 0.0
W1 relative error: 1.8336562786695002e-08
W2 relative error: 3.201560569143183e-10
b1 relative error: 9.828315204644842e-09
b2 relative error: 4.329134954569865e-10
Running numeric gradient check with reg = 0.7
W1 relative error: 2.5279152310200606e-07
W2 relative error: 2.8508510893102143e-08
b1 relative error: 1.564679947504764e-08
b2 relative error: 9.089617896905665e-10

```

## 1.7 Solver

We will now use the `cs231n Solver` class to train these networks. Familiarize yourself with the API in `cs231n/solver.py`. After you have done so, declare an instance of a `TwoLayerNet` with 200 units and then train it with the Solver. Choose parameters so that your validation accuracy is at least 50%.

```

In [10]: model = TwoLayerNet()
         solver = None

# ===== #
# YOUR CODE HERE:
#   Declare an instance of a TwoLayerNet and then train
#   it with the Solver. Choose hyperparameters so that your validation
#   accuracy is at least 50%. We won't have you optimize this further
#   since you did it in the previous notebook.
#
# ===== #

solver = Solver(model, data,
                update_rule='sgd',
                optim_config={
                    'learning_rate': 1e-3,
                },
                lr_decay=0.95,
                num_epochs=10, batch_size=100,
                print_every=100)
solver.train()

# ===== #
# END YOUR CODE HERE
# ===== #

```

```
(Iteration 1 / 4900) loss: 2.305584
```


(Epoch 0 / 10) train acc: 0.132000; val\_acc: 0.139000  
(Iteration 101 / 4900) loss: 1.793496  
(Iteration 201 / 4900) loss: 1.878698  
(Iteration 301 / 4900) loss: 1.528567  
(Iteration 401 / 4900) loss: 1.530954  
(Epoch 1 / 10) train acc: 0.407000; val\_acc: 0.447000  
(Iteration 501 / 4900) loss: 1.547538  
(Iteration 601 / 4900) loss: 1.706355  
(Iteration 701 / 4900) loss: 1.443473  
(Iteration 801 / 4900) loss: 1.463142  
(Iteration 901 / 4900) loss: 1.667293  
(Epoch 2 / 10) train acc: 0.488000; val\_acc: 0.484000  
(Iteration 1001 / 4900) loss: 1.525887  
(Iteration 1101 / 4900) loss: 1.388733  
(Iteration 1201 / 4900) loss: 1.466066  
(Iteration 1301 / 4900) loss: 1.451888  
(Iteration 1401 / 4900) loss: 1.549298  
(Epoch 3 / 10) train acc: 0.503000; val\_acc: 0.457000  
(Iteration 1501 / 4900) loss: 1.264786  
(Iteration 1601 / 4900) loss: 1.334163  
(Iteration 1701 / 4900) loss: 1.365588  
(Iteration 1801 / 4900) loss: 1.504482  
(Iteration 1901 / 4900) loss: 1.315015  
(Epoch 4 / 10) train acc: 0.542000; val\_acc: 0.493000  
(Iteration 2001 / 4900) loss: 1.150817  
(Iteration 2101 / 4900) loss: 1.446539  
(Iteration 2201 / 4900) loss: 1.187517  
(Iteration 2301 / 4900) loss: 1.394844  
(Iteration 2401 / 4900) loss: 1.197267  
(Epoch 5 / 10) train acc: 0.569000; val\_acc: 0.492000  
(Iteration 2501 / 4900) loss: 1.309897  
(Iteration 2601 / 4900) loss: 1.293208  
(Iteration 2701 / 4900) loss: 1.386833  
(Iteration 2801 / 4900) loss: 1.212871  
(Iteration 2901 / 4900) loss: 1.505262  
(Epoch 6 / 10) train acc: 0.584000; val\_acc: 0.505000  
(Iteration 3001 / 4900) loss: 1.118547  
(Iteration 3101 / 4900) loss: 1.182619  
(Iteration 3201 / 4900) loss: 1.307901  
(Iteration 3301 / 4900) loss: 1.082882  
(Iteration 3401 / 4900) loss: 1.152116  
(Epoch 7 / 10) train acc: 0.586000; val\_acc: 0.513000  
(Iteration 3501 / 4900) loss: 1.400135  
(Iteration 3601 / 4900) loss: 1.301806  
(Iteration 3701 / 4900) loss: 1.024732  
(Iteration 3801 / 4900) loss: 1.151949  
(Iteration 3901 / 4900) loss: 1.169311  
(Epoch 8 / 10) train acc: 0.574000; val\_acc: 0.479000

```
(Iteration 4001 / 4900) loss: 1.033320
(Iteration 4101 / 4900) loss: 1.285716
(Iteration 4201 / 4900) loss: 1.098872
(Iteration 4301 / 4900) loss: 1.186996
(Iteration 4401 / 4900) loss: 1.081218
(Epoch 9 / 10) train acc: 0.597000; val_acc: 0.507000
(Iteration 4501 / 4900) loss: 1.200677
(Iteration 4601 / 4900) loss: 1.038547
(Iteration 4701 / 4900) loss: 1.442947
(Iteration 4801 / 4900) loss: 1.136319
(Epoch 10 / 10) train acc: 0.620000; val_acc: 0.498000
```

In [11]: *# Run this cell to visualize training loss and train / val accuracy*

```
plt.subplot(2, 1, 1)
plt.title('Training loss')
plt.plot(solver.loss_history, 'o')
plt.xlabel('Iteration')

plt.subplot(2, 1, 2)
plt.title('Accuracy')
plt.plot(solver.train_acc_history, '-o', label='train')
plt.plot(solver.val_acc_history, '-o', label='val')
plt.plot([0.5] * len(solver.val_acc_history), 'k--')
plt.xlabel('Epoch')
plt.legend(loc='lower right')
plt.gcf().set_size_inches(15, 12)
plt.show()
```



## 1.8 Multilayer Neural Network

Now, we implement a multi-layer neural network.

Read through the `FullyConnectedNet` class in the file `nnd1/fc_net.py`.

Implement the initialization, the forward pass, and the backward pass. There will be lines for batchnorm and dropout layers and caches; ignore these all for now. That'll be in assignment #4.

```
In [12]: N, D, H1, H2, C = 2, 15, 20, 30, 10
         X = np.random.randn(N, D)
         y = np.random.randint(C, size=(N,))

         for reg in [0, 3.14]:
             print('Running check with reg = {}'.format(reg))
             model = FullyConnectedNet([H1, H2], input_dim=D, num_classes=C,
                                       reg=reg, weight_scale=5e-2, dtype=np.float64)

             loss, grads = model.loss(X, y)
             print('Initial loss: {}'.format(loss))
```

```

for name in sorted(grads):
    f = lambda _: model.loss(X, y)[0]
    grad_num = eval_numerical_gradient(f, model.params[name], verbose=False, h=1e-5)
    print('{} relative error: {}'.format(name, rel_error(grad_num, grads[name])))

```

Running check with reg = 0

Initial loss: 2.303593162473671

W1 relative error: 2.336046789069449e-06

W2 relative error: 2.8907377479669035e-07

W3 relative error: 8.530055965775117e-08

b1 relative error: 2.1196474895415304e-07

b2 relative error: 5.206315899482177e-09

b3 relative error: 1.3211502756366184e-10

Running check with reg = 3.14

Initial loss: 6.667576460326028

W1 relative error: 1.266501424135116e-08

W2 relative error: 2.429748660718304e-07

W3 relative error: 1.9289959633719516e-07

b1 relative error: 2.877261542611307e-08

b2 relative error: 3.741208373313691e-09

b3 relative error: 1.5668469055220386e-10

In [13]: *# Use the three layer neural network to overfit a small dataset.*

```

num_train = 50
small_data = {
    'X_train': data['X_train'][:num_train],
    'y_train': data['y_train'][:num_train],
    'X_val': data['X_val'],
    'y_val': data['y_val'],
}

#### !!!!!
# Play around with the weight_scale and learning_rate so that you can overfit a small
# Your training accuracy should be 1.0 to receive full credit on this part.
weight_scale = 1e-2
learning_rate = 1e-2

model = FullyConnectedNet([100, 100],
                           weight_scale=weight_scale, dtype=np.float64)
solver = Solver(model, small_data,
                print_every=10, num_epochs=20, batch_size=25,
                update_rule='sgd',
                optim_config={
                    'learning_rate': learning_rate,
                })

```

```

    )
    solver.train()

    plt.plot(solver.loss_history, 'o')
    plt.title('Training loss history')
    plt.xlabel('Iteration')
    plt.ylabel('Training loss')
    plt.show()

(Iteration 1 / 40) loss: 2.266194
(Epoch 0 / 20) train acc: 0.160000; val_acc: 0.087000
(Epoch 1 / 20) train acc: 0.380000; val_acc: 0.112000
(Epoch 2 / 20) train acc: 0.360000; val_acc: 0.140000
(Epoch 3 / 20) train acc: 0.560000; val_acc: 0.152000
(Epoch 4 / 20) train acc: 0.620000; val_acc: 0.151000
(Epoch 5 / 20) train acc: 0.680000; val_acc: 0.186000
(Iteration 11 / 40) loss: 0.997419
(Epoch 6 / 20) train acc: 0.680000; val_acc: 0.152000
(Epoch 7 / 20) train acc: 0.780000; val_acc: 0.155000
(Epoch 8 / 20) train acc: 0.780000; val_acc: 0.166000
(Epoch 9 / 20) train acc: 0.920000; val_acc: 0.179000
(Epoch 10 / 20) train acc: 0.900000; val_acc: 0.177000
(Iteration 21 / 40) loss: 0.244519
(Epoch 11 / 20) train acc: 0.980000; val_acc: 0.174000
(Epoch 12 / 20) train acc: 0.940000; val_acc: 0.169000
(Epoch 13 / 20) train acc: 0.960000; val_acc: 0.177000
(Epoch 14 / 20) train acc: 0.980000; val_acc: 0.169000
(Epoch 15 / 20) train acc: 0.980000; val_acc: 0.178000
(Iteration 31 / 40) loss: 0.067768
(Epoch 16 / 20) train acc: 0.980000; val_acc: 0.163000
(Epoch 17 / 20) train acc: 1.000000; val_acc: 0.170000
(Epoch 18 / 20) train acc: 1.000000; val_acc: 0.168000
(Epoch 19 / 20) train acc: 1.000000; val_acc: 0.174000
(Epoch 20 / 20) train acc: 1.000000; val_acc: 0.161000

```





```

1 import numpy as np
2
3 from .layers import *
4 from .layer_utils import *
5
6 """
7 This code was originally written for CS 231n at Stanford University
8 (cs231n.stanford.edu). It has been modified in various areas for use in the
9 ECE 239AS class at UCLA. This includes the descriptions of what code to
10 implement as well as some slight potential changes in variable names to be
11 consistent with class nomenclature. We thank Justin Johnson & Serena Yeung
12 for
13 permission to use this code. To see the original version, please visit
14 cs231n.stanford.edu.
15 """
16 class TwoLayerNet(object):
17     """
18     A two-layer fully-connected neural network with ReLU nonlinearity and
19     softmax loss that uses a modular layer design. We assume an input dimension
20     of D, a hidden dimension of H, and perform classification over C classes.
21
22     The architecture should be affine - relu - affine - softmax.
23
24     Note that this class does not implement gradient descent; instead, it
25     will interact with a separate Solver object that is responsible for running
26     optimization.
27
28     The learnable parameters of the model are stored in the dictionary
29     self.params that maps parameter names to numpy arrays.
30     """
31
32     def __init__(self, input_dim=3*32*32, hidden_dims=100, num_classes=10,
33                 dropout=0, weight_scale=1e-3, reg=0.0):
34         """
35         Initialize a new network.
36
37         Inputs:
38         - input_dim: An integer giving the size of the input
39         - hidden_dims: An integer giving the size of the hidden layer
40         - num_classes: An integer giving the number of classes to classify
41         - dropout: Scalar between 0 and 1 giving dropout strength.
42         - weight_scale: Scalar giving the standard deviation for random
43           initialization of the weights.
44         - reg: Scalar giving L2 regularization strength.
45         """
46         self.params = {}
47         self.reg = reg
48
49         # ===== #
50         # YOUR CODE HERE:
51         # Initialize W1, W2, b1, and b2. Store these as self.params['W1'],
52         # self.params['W2'], self.params['b1'] and self.params['b2']. The
53         # biases are initialized to zero and the weights are initialized
54         # so that each parameter has mean 0 and standard deviation
55         weight_scale.
56         # The dimensions of W1 should be (input_dim, hidden_dim) and the
57         # dimensions of W2 should be (hidden_dims, num_classes)
58         # ===== #

```

```

59     self.params['W1'] = weight_scale * np.random.randn(input_dim,
hidden_dims)
60     self.params['b1'] = np.zeros(hidden_dims)
61     self.params['W2'] = weight_scale * np.random.randn(hidden_dims,
num_classes)
62     self.params['b2'] = np.zeros(num_classes)
63
64     # ===== #
65     # END YOUR CODE HERE
66     # ===== #
67
68 def loss(self, X, y=None):
69     """
70     Compute loss and gradient for a minibatch of data.
71
72     Inputs:
73     - X: Array of input data of shape (N, d_1, ..., d_k)
74     - y: Array of labels, of shape (N,). y[i] gives the label for X[i].
75
76     Returns:
77     If y is None, then run a test-time forward pass of the model and return:
78     - scores: Array of shape (N, C) giving classification scores, where
79       scores[i, c] is the classification score for X[i] and class c.
80
81     If y is not None, then run a training-time forward and backward pass and
82     return a tuple of:
83     - loss: Scalar value giving the loss
84     - grads: Dictionary with the same keys as self.params, mapping parameter
85       names to gradients of the loss with respect to those parameters.
86     """
87     scores = None
88
89     # ===== #
90     # YOUR CODE HERE:
91     #   Implement the forward pass of the two-layer neural network. Store
92     #   the class scores as the variable 'scores'. Be sure to use the layers
93     #   you prior implemented.
94     # ===== #
95
96     out, l1_cache = affine_relu_forward(X, self.params['W1'],
self.params['b1'])
97     scores, l2_cache = affine_forward(out, self.params['W2'],
self.params['b2'])
98
99     # ===== #
100    # END YOUR CODE HERE
101    # ===== #
102
103    # If y is None then we are in test mode so just return scores
104    if y is None:
105        return scores
106
107    loss, grads = 0, {}
108    # ===== #
109    # YOUR CODE HERE:
110    #   Implement the backward pass of the two-layer neural net. Store
111    #   the loss as the variable 'loss' and store the gradients in the
112    #   'grads' dictionary. For the grads dictionary, grads['W1'] holds
113    #   the gradient for W1, grads['b1'] holds the gradient for b1, etc.
114    #   i.e., grads[k] holds the gradient for self.params[k].

```

```

115 #
116 # Add L2 regularization, where there is an added cost  $0.5 * \text{self.reg} * W^2$ 
117 # for each W. Be sure to include the 0.5 multiplying factor to
118 # match our implementation.
119 #
120 # And be sure to use the layers you prior implemented.
121 # ===== #
122
123 loss, dL = softmax_loss(scores, y)
124 reg_loss = 0.5 * self.reg * (np.sum(self.params['W1']**2) +
np.sum(self.params['W2']**2))
125 loss += reg_loss
126
127 dh1, grads['W2'], grads['b2'] = affine_backward(dL, l2_cache)
128 dx, grads['W1'], grads['b1'] = affine_relu_backward(dh1, l1_cache)
129
130 grads['W1'] += self.reg * self.params['W1']
131 grads['W2'] += self.reg * self.params['W2']
132
133 # ===== #
134 # END YOUR CODE HERE
135 # ===== #
136
137 return loss, grads
138
139
140 class FullyConnectedNet(object):
141     """
142     A fully-connected neural network with an arbitrary number of hidden layers,
143     ReLU nonlinearities, and a softmax loss function. This will also implement
144     dropout and batch normalization as options. For a network with L layers,
145     the architecture will be
146
147     {affine - [batch norm] - relu - [dropout]} x (L - 1) - affine - softmax
148
149     where batch normalization and dropout are optional, and the {...} block is
150     repeated L - 1 times.
151
152     Similar to the TwoLayerNet above, learnable parameters are stored in the
153     self.params dictionary and will be learned using the Solver class.
154     """
155
156     def __init__(self, hidden_dims, input_dim=3*32*32, num_classes=10,
157                 dropout=0, use_batchnorm=False, reg=0.0,
158                 weight_scale=1e-2, dtype=np.float32, seed=None):
159         """
160         Initialize a new FullyConnectedNet.
161
162         Inputs:
163         - hidden_dims: A list of integers giving the size of each hidden layer.
164         - input_dim: An integer giving the size of the input.
165         - num_classes: An integer giving the number of classes to classify.
166         - dropout: Scalar between 0 and 1 giving dropout strength. If dropout=0
then
167     the network should not use dropout at all.
168     - use_batchnorm: Whether or not the network should use batch
normalization.
169     - reg: Scalar giving L2 regularization strength.
170     - weight_scale: Scalar giving the standard deviation for random
171     initialization of the weights.

```

```

172     - dtype: A numpy datatype object; all computations will be performed
using
173     this datatype. float32 is faster but less accurate, so you should use
174     float64 for numeric gradient checking.
175     - seed: If not None, then pass this random seed to the dropout layers.
This
176     will make the dropout layers deterministic so we can gradient check the
177     model.
178     """
179     self.use_batchnorm = use_batchnorm
180     self.use_dropout = dropout > 0
181     self.reg = reg
182     self.num_layers = 1 + len(hidden_dims)
183     self.dtype = dtype
184     self.params = {}
185
186     # ===== #
187     # YOUR CODE HERE:
188     # Initialize all parameters of the network in the self.params
dictionary.
189     # The weights and biases of layer 1 are W1 and b1; and in general the
190     # weights and biases of layer i are Wi and bi. The
191     # biases are initialized to zero and the weights are initialized
192     # so that each parameter has mean 0 and standard deviation
weight_scale.
193     # ===== #
194
195     dims = hidden_dims
196     dims.insert(0, input_dim)
197     dims.append(num_classes)
198
199     for i in range(len(dims) - 1):
200         self.params['W{}'.format(i+1)] = weight_scale *
np.random.randn(dims[i], dims[i+1])
201         self.params['b{}'.format(i+1)] = np.zeros(dims[i+1])
202
203     # ===== #
204     # END YOUR CODE HERE
205     # ===== #
206
207     # When using dropout we need to pass a dropout_param dictionary to each
208     # dropout layer so that the layer knows the dropout probability and the
mode
209     # (train / test). You can pass the same dropout_param to each dropout
layer.
210     self.dropout_param = {}
211     if self.use_dropout:
212         self.dropout_param = {'mode': 'train', 'p': dropout}
213         if seed is not None:
214             self.dropout_param['seed'] = seed
215
216     # With batch normalization we need to keep track of running means and
217     # variances, so we need to pass a special bn_param object to each batch
218     # normalization layer. You should pass self.bn_params[0] to the forward
pass
219     # of the first batch normalization layer, self.bn_params[1] to the
forward
220     # pass of the second batch normalization layer, etc.
221     self.bn_params = []
222     if self.use_batchnorm:

```

```

223     self.bn_params = [{'mode': 'train'} for i in np.arange(self.num_layers
- 1)]
224
225     # Cast all parameters to the correct datatype
226     for k, v in self.params.items():
227         self.params[k] = v.astype(dtype)
228
229
230 def loss(self, X, y=None):
231     """
232     Compute loss and gradient for the fully-connected net.
233
234     Input / output: Same as TwoLayerNet above.
235     """
236     X = X.astype(self.dtype)
237     mode = 'test' if y is None else 'train'
238
239     # Set train/test mode for batchnorm params and dropout param since they
240     # behave differently during training and testing.
241     if self.dropout_param is not None:
242         self.dropout_param['mode'] = mode
243     if self.use_batchnorm:
244         for bn_param in self.bn_params:
245             bn_param[mode] = mode
246
247     scores = None
248
249     # ===== #
250     # YOUR CODE HERE:
251     # Implement the forward pass of the FC net and store the output
252     # scores as the variable "scores".
253     # ===== #
254
255     caches = []
256     x = X
257     for i in range(1, self.num_layers + 1):
258         w, b = self.params['W{}'.format(i)], self.params['b{}'.format(i)]
259         if (i == self.num_layers):
260             scores, cache = affine_forward(x, w, b)
261         else:
262             x, cache = affine_relu_forward(x, w, b)
263         caches.append(cache)
264
265     # ===== #
266     # END YOUR CODE HERE
267     # ===== #
268
269     # If test mode return early
270     if mode == 'test':
271         return scores
272
273     loss, grads = 0.0, {}
274     # ===== #
275     # YOUR CODE HERE:
276     # Implement the backwards pass of the FC net and store the gradients
277     # in the grads dict, so that grads[k] is the gradient of self.params[k]
278     # Be sure your L2 regularization includes a 0.5 factor.
279     # ===== #
280
281     sm_loss, dout = softmax_loss(scores, y)

```

```

282     reg_loss = 0.0
283     for i in range(1, self.num_layers + 1):
284         reg_loss += 0.5 * self.reg * (np.sum(self.params['W{}'.format(i)]**2))
285     loss = sm_loss + reg_loss
286
287     for i in range(self.num_layers, 0, -1):
288         cur_w = 'W{}'.format(i)
289         cur_b = 'b{}'.format(i)
290         if (i == self.num_layers):
291             dout, grads[cur_w], grads[cur_b] = affine_backward(dout,
292 caches.pop())
293         else:
294             dout, grads[cur_w], grads[cur_b] = affine_relu_backward(dout,
295 caches.pop())
296             grads[cur_w] += self.reg * self.params[cur_w]
297
298     # ===== #
299     # END YOUR CODE HERE
300     # ===== #
301     return loss, grads

```



```

1 import numpy as np
2 import pdb
3
4 """
5 This code was originally written for CS 231n at Stanford University
6 (cs231n.stanford.edu). It has been modified in various areas for use in the
7 ECE 239AS class at UCLA. This includes the descriptions of what code to
8 implement as well as some slight potential changes in variable names to be
9 consistent with class nomenclature. We thank Justin Johnson & Serena Yeung
10 for
11 permission to use this code. To see the original version, please visit
12 cs231n.stanford.edu.
13 """
14
15 def affine_forward(x, w, b):
16     """
17     Computes the forward pass for an affine (fully-connected) layer.
18
19     The input x has shape (N, d_1, ..., d_k) and contains a minibatch of N
20     examples, where each example x[i] has shape (d_1, ..., d_k). We will
21     reshape each input into a vector of dimension D = d_1 * ... * d_k, and
22     then transform it to an output vector of dimension M.
23
24     Inputs:
25     - x: A numpy array containing input data, of shape (N, d_1, ..., d_k)
26     - w: A numpy array of weights, of shape (D, M)
27     - b: A numpy array of biases, of shape (M,)
28
29     Returns a tuple of:
30     - out: output, of shape (N, M)
31     - cache: (x, w, b)
32     """
33
34     # ===== #
35     # YOUR CODE HERE:
36     # Calculate the output of the forward pass. Notice the dimensions
37     # of w are D x M, which is the transpose of what we did in earlier
38     # assignments.
39     # ===== #
40
41     X = x.reshape(x.shape[0], -1)
42     out = X.dot(w) + b
43
44     # ===== #
45     # END YOUR CODE HERE
46     # ===== #
47
48     cache = (x, w, b)
49     return out, cache
50
51
52 def affine_backward(dout, cache):
53     """
54     Computes the backward pass for an affine layer.
55
56     Inputs:
57     - dout: Upstream derivative, of shape (N, M)
58     - cache: Tuple of:
59       - x: Input data, of shape (N, d_1, ..., d_k)

```

```

60     - w: Weights, of shape (D, M)
61
62     Returns a tuple of:
63     - dx: Gradient with respect to x, of shape (N, d1, ..., d_k)
64     - dw: Gradient with respect to w, of shape (D, M)
65     - db: Gradient with respect to b, of shape (M,)
66     """
67     x, w, b = cache
68     dx, dw, db = None, None, None
69
70     # ===== #
71     # YOUR CODE HERE:
72     #     Calculate the gradients for the backward pass.
73     # ===== #
74
75     # dout is N x M
76     # dx should be N x d1 x ... x dk; it relates to dout through multiplication
with w, which is D x M
77     # dw should be D x M; it relates to dout through multiplication with x,
which is N x D after reshaping
78     # db should be M; it is just the sum over dout examples
79
80     dx = dout.dot(w.T).reshape(x.shape)
81     dw = x.reshape(x.shape[0], -1).T.dot(dout)
82     db = dout.sum(axis=0)
83
84     # ===== #
85     # END YOUR CODE HERE
86     # ===== #
87
88     return dx, dw, db
89
90 def relu_forward(x):
91     """
92     Computes the forward pass for a layer of rectified linear units (ReLUs).
93
94     Input:
95     - x: Inputs, of any shape
96
97     Returns a tuple of:
98     - out: Output, of the same shape as x
99     - cache: x
100    """
101    # ===== #
102    # YOUR CODE HERE:
103    #     Implement the ReLU forward pass.
104    # ===== #
105
106    out = x * (x > 0)
107
108    # ===== #
109    # END YOUR CODE HERE
110    # ===== #
111
112    cache = x
113    return out, cache
114
115
116 def relu_backward(dout, cache):
117     """

```

```

118 Computes the backward pass for a layer of rectified linear units (ReLUs).
119
120 Input:
121 - dout: Upstream derivatives, of any shape
122 - cache: Input x, of same shape as dout
123
124 Returns:
125 - dx: Gradient with respect to x
126 """
127 x = cache
128
129 # ===== #
130 # YOUR CODE HERE:
131 #   Implement the ReLU backward pass
132 # ===== #
133
134 # ReLU directs linearly to those > 0
135
136 dx = dout * (cache > 0)
137
138 # ===== #
139 # END YOUR CODE HERE
140 # ===== #
141
142 return dx
143
144 def svm_loss(x, y):
145     """
146     Computes the loss and gradient using for multiclass SVM classification.
147
148     Inputs:
149     - x: Input data, of shape (N, C) where x[i, j] is the score for the jth
class
150     for the ith input.
151     - y: Vector of labels, of shape (N,) where y[i] is the label for x[i] and
152     0 <= y[i] < C
153
154     Returns a tuple of:
155     - loss: Scalar giving the loss
156     - dx: Gradient of the loss with respect to x
157     """
158     N = x.shape[0]
159     correct_class_scores = x[np.arange(N), y]
160     margins = np.maximum(0, x - correct_class_scores[:, np.newaxis] + 1.0)
161     margins[np.arange(N), y] = 0
162     loss = np.sum(margins) / N
163     num_pos = np.sum(margins > 0, axis=1)
164     dx = np.zeros_like(x)
165     dx[margins > 0] = 1
166     dx[np.arange(N), y] -= num_pos
167     dx /= N
168     return loss, dx
169
170
171 def softmax_loss(x, y):
172     """
173     Computes the loss and gradient for softmax classification.
174
175     Inputs:

```

```

176     - x: Input data, of shape (N, C) where x[i, j] is the score for the jth
class
177     for the ith input.
178     - y: Vector of labels, of shape (N,) where y[i] is the label for x[i] and
179         0 <= y[i] < C
180
181     Returns a tuple of:
182     - loss: Scalar giving the loss
183     - dx: Gradient of the loss with respect to x
184     """
185
186     probs = np.exp(x - np.max(x, axis=1, keepdims=True))
187     probs /= np.sum(probs, axis=1, keepdims=True)
188     N = x.shape[0]
189     loss = -np.sum(np.log(probs[np.arange(N), y])) / N
190     dx = probs.copy()
191     dx[np.arange(N), y] -= 1
192     dx /= N
193     return loss, dx
194

```