

```

1 import numpy as np
2 import pdb
3
4 """
5 This code was originally written for CS 231n at Stanford University
6 (cs231n.stanford.edu). It has been modified in various areas for use in the
7 ECE 239AS class at UCLA. This includes the descriptions of what code to
8 implement as well as some slight potential changes in variable names to be
9 consistent with class nomenclature. We thank Justin Johnson & Serena Yeung
10 for
11 permission to use this code. To see the original version, please visit
12 cs231n.stanford.edu.
13 """
14
15 def affine_forward(x, w, b):
16     """
17     Computes the forward pass for an affine (fully-connected) layer.
18
19     The input x has shape (N, d_1, ..., d_k) and contains a minibatch of N
20     examples, where each example x[i] has shape (d_1, ..., d_k). We will
21     reshape each input into a vector of dimension D = d_1 * ... * d_k, and
22     then transform it to an output vector of dimension M.
23
24     Inputs:
25     - x: A numpy array containing input data, of shape (N, d_1, ..., d_k)
26     - w: A numpy array of weights, of shape (D, M)
27     - b: A numpy array of biases, of shape (M,)
28
29     Returns a tuple of:
30     - out: output, of shape (N, M)
31     - cache: (x, w, b)
32     """
33
34     # ===== #
35     # YOUR CODE HERE:
36     # Calculate the output of the forward pass. Notice the dimensions
37     # of w are D x M, which is the transpose of what we did in earlier
38     # assignments.
39     # ===== #
40
41     X = x.reshape(x.shape[0], -1)
42     out = X.dot(w) + b
43
44     # ===== #
45     # END YOUR CODE HERE
46     # ===== #
47
48     cache = (x, w, b)
49     return out, cache
50
51
52 def affine_backward(dout, cache):
53     """
54     Computes the backward pass for an affine layer.
55
56     Inputs:
57     - dout: Upstream derivative, of shape (N, M)
58     - cache: Tuple of:
59       - x: Input data, of shape (N, d_1, ..., d_k)

```

```

60     - w: Weights, of shape (D, M)
61
62     Returns a tuple of:
63     - dx: Gradient with respect to x, of shape (N, d1, ..., d_k)
64     - dw: Gradient with respect to w, of shape (D, M)
65     - db: Gradient with respect to b, of shape (M,)
66     """
67     x, w, b = cache
68     dx, dw, db = None, None, None
69
70     # ===== #
71     # YOUR CODE HERE:
72     #     Calculate the gradients for the backward pass.
73     # ===== #
74
75     # dout is N x M
76     # dx should be N x d1 x ... x dk; it relates to dout through multiplication
with w, which is D x M
77     # dw should be D x M; it relates to dout through multiplication with x,
which is N x D after reshaping
78     # db should be M; it is just the sum over dout examples
79
80     dx = dout.dot(w.T).reshape(x.shape)
81     dw = x.reshape(x.shape[0], -1).T.dot(dout)
82     db = dout.sum(axis=0)
83
84     # ===== #
85     # END YOUR CODE HERE
86     # ===== #
87
88     return dx, dw, db
89
90 def relu_forward(x):
91     """
92     Computes the forward pass for a layer of rectified linear units (ReLUs).
93
94     Input:
95     - x: Inputs, of any shape
96
97     Returns a tuple of:
98     - out: Output, of the same shape as x
99     - cache: x
100     """
101     # ===== #
102     # YOUR CODE HERE:
103     #     Implement the ReLU forward pass.
104     # ===== #
105
106     out = x * (x > 0)
107
108     # ===== #
109     # END YOUR CODE HERE
110     # ===== #
111
112     cache = x
113     return out, cache
114
115
116 def relu_backward(dout, cache):
117     """

```

```

118 Computes the backward pass for a layer of rectified linear units (ReLUs).
119
120 Input:
121 - dout: Upstream derivatives, of any shape
122 - cache: Input x, of same shape as dout
123
124 Returns:
125 - dx: Gradient with respect to x
126 """
127 x = cache
128
129 # ===== #
130 # YOUR CODE HERE:
131 #   Implement the ReLU backward pass
132 # ===== #
133
134 # ReLU directs linearly to those > 0
135
136 dx = dout * (cache > 0)
137
138 # ===== #
139 # END YOUR CODE HERE
140 # ===== #
141
142 return dx
143
144 def svm_loss(x, y):
145     """
146     Computes the loss and gradient using for multiclass SVM classification.
147
148     Inputs:
149     - x: Input data, of shape (N, C) where x[i, j] is the score for the jth
class
150     for the ith input.
151     - y: Vector of labels, of shape (N,) where y[i] is the label for x[i] and
152     0 <= y[i] < C
153
154     Returns a tuple of:
155     - loss: Scalar giving the loss
156     - dx: Gradient of the loss with respect to x
157     """
158     N = x.shape[0]
159     correct_class_scores = x[np.arange(N), y]
160     margins = np.maximum(0, x - correct_class_scores[:, np.newaxis] + 1.0)
161     margins[np.arange(N), y] = 0
162     loss = np.sum(margins) / N
163     num_pos = np.sum(margins > 0, axis=1)
164     dx = np.zeros_like(x)
165     dx[margins > 0] = 1
166     dx[np.arange(N), y] -= num_pos
167     dx /= N
168     return loss, dx
169
170
171 def softmax_loss(x, y):
172     """
173     Computes the loss and gradient for softmax classification.
174
175     Inputs:

```

```
176     - x: Input data, of shape (N, C) where x[i, j] is the score for the jth
class
177     for the ith input.
178     - y: Vector of labels, of shape (N,) where y[i] is the label for x[i] and
179         0 <= y[i] < C
180
181     Returns a tuple of:
182     - loss: Scalar giving the loss
183     - dx: Gradient of the loss with respect to x
184     """
185
186     probs = np.exp(x - np.max(x, axis=1, keepdims=True))
187     probs /= np.sum(probs, axis=1, keepdims=True)
188     N = x.shape[0]
189     loss = -np.sum(np.log(probs[np.arange(N), y])) / N
190     dx = probs.copy()
191     dx[np.arange(N), y] -= 1
192     dx /= N
193     return loss, dx
194
```