

```

1 import numpy as np
2 import matplotlib.pyplot as plt
3
4 """
5 This code was originally written for CS 231n at Stanford University
6 (cs231n.stanford.edu). It has been modified in various areas for use in the
7 ECE 239AS class at UCLA. This includes the descriptions of what code to
8 implement as well as some slight potential changes in variable names to be
9 consistent with class nomenclature. We thank Justin Johnson & Serena Yeung
10 for
11 permission to use this code. To see the original version, please visit
12 cs231n.stanford.edu.
13 """
14
15 class TwoLayerNet(object):
16     """
17     A two-layer fully-connected neural network. The net has an input dimension
18     of
19     N, a hidden layer dimension of H, and performs classification over C
20     classes.
21     We train the network with a softmax loss function and L2 regularization on the
22     weight matrices. The network uses a ReLU nonlinearity after the first fully
23     connected layer.
24
25     In other words, the network has the following architecture:
26
27     input - fully connected layer - ReLU - fully connected layer - softmax
28
29     The outputs of the second fully-connected layer are the scores for each
30     class.
31     """
32
33     def __init__(self, input_size, hidden_size, output_size, std=1e-4):
34         """
35         Initialize the model. Weights are initialized to small random values and
36         biases are initialized to zero. Weights and biases are stored in the
37         variable self.params, which is a dictionary with the following keys:
38
39         W1: First layer weights; has shape (H, D)
40         b1: First layer biases; has shape (H,)
41         W2: Second layer weights; has shape (C, H)
42         b2: Second layer biases; has shape (C,)
43
44         Inputs:
45         - input_size: The dimension D of the input data.
46         - hidden_size: The number of neurons H in the hidden layer.
47         - output_size: The number of classes C.
48         """
49         self.params = {}
50         self.params['W1'] = std * np.random.randn(hidden_size, input_size)
51         self.params['b1'] = np.zeros(hidden_size)
52         self.params['W2'] = std * np.random.randn(output_size, hidden_size)
53         self.params['b2'] = np.zeros(output_size)
54
55     def loss(self, X, y=None, reg=0.0):
56         """
57         Compute the loss and gradients for a two layer fully connected neural
58         network.

```

```

56
57 Inputs:
58 - X: Input data of shape (N, D). Each X[i] is a training sample.
59 - y: Vector of training labels. y[i] is the label for X[i], and each y[i]
is
60 an integer in the range  $0 \leq y[i] < C$ . This parameter is optional; if
it
61 is not passed then we only return scores, and if it is passed then we
62 instead return the loss and gradients.
63 - reg: Regularization strength.
64
65 Returns:
66 If y is None, return a matrix scores of shape (N, C) where scores[i, c]
is
67 the score for class c on input X[i].
68
69 If y is not None, instead return a tuple of:
70 - loss: Loss (data loss and regularization loss) for this batch of
training
71 samples.
72 - grads: Dictionary mapping parameter names to gradients of those
parameters
73 with respect to the loss function; has the same keys as self.params.
74 """
75 # Unpack variables from the params dictionary
76 W1, b1 = self.params['W1'], self.params['b1']
77 W2, b2 = self.params['W2'], self.params['b2']
78 N, D = X.shape
79
80 # Compute the forward pass
81 scores = None
82
83 # ===== #
84 # YOUR CODE HERE:
85 # Calculate the output scores of the neural network. The result
86 # should be (N, C). As stated in the description for this class,
87 # there should not be a ReLU layer after the second FC layer.
88 # The output of the second FC layer is the output scores. Do not
89 # use a for loop in your implementation.
90 # ===== #
91
92 h1 = np.maximum(X.dot(W1.T) + b1, 0)
93 h2 = h1.dot(W2.T) + b2
94 scores = h2
95
96 # ===== #
97 # END YOUR CODE HERE
98 # ===== #
99
100 # If the targets are not given then jump out, we're done
101 if y is None:
102     return scores
103
104 # Compute the loss
105 loss = None
106
107 # ===== #
108 # YOUR CODE HERE:
109 # Calculate the loss of the neural network. This includes the
110 # softmax loss and the L2 regularization for W1 and W2. Store the

```

```

111 # total loss in the variable loss. Multiply the regularization
112 #   loss by 0.5 (in addition to the factor reg).
113 # ===== #
114
115     # scores is num_examples by num_classes
116     ea = np.exp(scores - np.max(scores))
117     sums = np.sum(ea, axis=1)
118     softmax = ea / sums[:, np.newaxis]
119     loss = np.mean(-np.log(softmax[np.arange(N), y]))
120     loss += 0.5 * reg * (np.sum(W1**2) + np.sum(W2**2))
121     # ===== #
122     # END YOUR CODE HERE
123     # ===== #
124
125     grads = {}
126
127     # ===== #
128     # YOUR CODE HERE:
129 #   Implement the backward pass. Compute the derivatives of the
130 #   weights and the biases. Store the results in the grads
131 #   dictionary. e.g., grads['W1'] should store the gradient for
132 #   W1, and be of the same size as W1.
133 # ===== #
134
135     grad_softmax = softmax
136     grad_softmax[np.arange(N), y] -= 1
137     grad_softmax /= N
138     grads['W2'] = grad_softmax.T.dot(h1) + reg * W2    # (C x N) * (N x H) =
(C x H)
139     grads['b2'] = np.sum(grad_softmax, axis=0)        # (C, )
140     # (H x C) * (C x N) = (H x N)
141     # (H x N) * (N x D) = (H x D)
142     ind = h1
143     ind[ind > 0] = 1
144     ind[ind <= 0] = 0
145     grads['W1'] = (W2.T.dot(grad_softmax.T) * ind.T).dot(X) + reg * W1
146     grads['b1'] = np.sum((W2.T.dot(grad_softmax.T) * ind.T).T, axis = 0)
147
148     # ===== #
149     # END YOUR CODE HERE
150     # ===== #
151
152     return loss, grads
153
154 def train(self, X, y, X_val, y_val,
155           learning_rate=1e-3, learning_rate_decay=0.95,
156           reg=1e-5, num_iters=100,
157           batch_size=200, verbose=False):
158     """
159     Train this neural network using stochastic gradient descent.
160
161     Inputs:
162     - X: A numpy array of shape (N, D) giving training data.
163     - y: A numpy array of shape (N,) giving training labels; y[i] = c means
that
164     X[i] has label c, where 0 <= c < C.
165     - X_val: A numpy array of shape (N_val, D) giving validation data.
166     - y_val: A numpy array of shape (N_val,) giving validation labels.
167     - learning_rate: Scalar giving learning rate for optimization.

```

```

168     - learning_rate_decay: Scalar giving factor used to decay the learning
rate
169         after each epoch.
170     - reg: Scalar giving regularization strength.
171     - num_iters: Number of steps to take when optimizing.
172     - batch_size: Number of training examples to use per step.
173     - verbose: boolean; if true print progress during optimization.
174     """
175     num_train = X.shape[0]
176     iterations_per_epoch = max(num_train / batch_size, 1)
177
178     # Use SGD to optimize the parameters in self.model
179     loss_history = []
180     train_acc_history = []
181     val_acc_history = []
182
183     for it in np.arange(num_iters):
184         X_batch = None
185         y_batch = None
186
187         # ===== #
188         # YOUR CODE HERE:
189         # Create a minibatch by sampling batch_size samples randomly.
190         # ===== #
191         indices = np.random.choice(num_train, batch_size)
192         X_batch = X[indices]
193         y_batch = y[indices]
194
195         # ===== #
196         # END YOUR CODE HERE
197         # ===== #
198
199         # Compute loss and gradients using the current minibatch
200         loss, grads = self.loss(X_batch, y=y_batch, reg=reg)
201         loss_history.append(loss)
202
203         # ===== #
204         # YOUR CODE HERE:
205         # Perform a gradient descent step using the minibatch to update
206         # all parameters (i.e., W1, W2, b1, and b2).
207         # ===== #
208
209         self.params['W1'] -= learning_rate * grads['W1']
210         self.params['b1'] -= learning_rate * grads['b1']
211         self.params['W2'] -= learning_rate * grads['W2']
212         self.params['b2'] -= learning_rate * grads['b2']
213
214         # ===== #
215         # END YOUR CODE HERE
216         # ===== #
217
218         if verbose and it % 100 == 0:
219             print('iteration {} / {}: loss {}'.format(it, num_iters, loss))
220
221         # Every epoch, check train and val accuracy and decay learning rate.
222         if it % iterations_per_epoch == 0:
223             # Check accuracy
224             train_acc = (self.predict(X_batch) == y_batch).mean()
225             val_acc = (self.predict(X_val) == y_val).mean()
226             train_acc_history.append(train_acc)

```

```

227         val_acc_history.append(val_acc)
228
229         # Decay learning rate
230         learning_rate *= learning_rate_decay
231
232     return {
233         'loss_history': loss_history,
234         'train_acc_history': train_acc_history,
235         'val_acc_history': val_acc_history,
236     }
237
238     def predict(self, X):
239         """
240         Use the trained weights of this two-layer network to predict labels for
241         data points. For each data point we predict scores for each of the C
242         classes, and assign each data point to the class with the highest score.
243
244         Inputs:
245         - X: A numpy array of shape (N, D) giving N D-dimensional data points to
246             classify.
247
248         Returns:
249         - y_pred: A numpy array of shape (N,) giving predicted labels for each of
250             the elements of X. For all i, y_pred[i] = c means that X[i] is
251             predicted
252             to have class c, where 0 <= c < C.
253         """
254         y_pred = None
255
256         # ===== #
257         # YOUR CODE HERE:
258         # Predict the class given the input data.
259         # ===== #
260
261         scores = np.maximum(X.dot(self.params['W1']).T + self.params['b1'],
262                             0).dot(self.params['W2']).T + self.params['b2']
263         y_pred = np.argmax(scores, axis=1)
264
265         # ===== #
266         # END YOUR CODE HERE
267         # ===== #
268
269         return y_pred

```