

```

1 import numpy as np
2 import pdb
3 import copy
4
5 from .layers import *
6 from .layer_utils import *
7
8 """
9 This code was originally written for CS 231n at Stanford University
10 (cs231n.stanford.edu). It has been modified in various areas for use in the
11 ECE 239AS class at UCLA. This includes the descriptions of what code to
12 implement as well as some slight potential changes in variable names to be
13 consistent with class nomenclature. We thank Justin Johnson & Serena Yeung
14 for
15 permission to use this code. To see the original version, please visit
16 cs231n.stanford.edu.
17 """
18 class TwoLayerNet(object):
19     """
20     A two-layer fully-connected neural network with ReLU nonlinearity and
21     softmax loss that uses a modular layer design. We assume an input
22     dimension
23     of D, a hidden dimension of H, and perform classification over C classes.
24
25     The architecture should be affine - relu - affine - softmax.
26
27     Note that this class does not implement gradient descent; instead, it
28     will interact with a separate Solver object that is responsible for
29     running
30     optimization.
31
32     The learnable parameters of the model are stored in the dictionary
33     self.params that maps parameter names to numpy arrays.
34     """
35     def __init__(self, input_dim=3*32*32, hidden_dims=100, num_classes=10,
36                 dropout=0, weight_scale=1e-3, reg=0.0):
37         """
38         Initialize a new network.
39
40         Inputs:
41         - input_dim: An integer giving the size of the input
42         - hidden_dims: An integer giving the size of the hidden layer
43         - num_classes: An integer giving the number of classes to classify
44         - dropout: Scalar between 0 and 1 giving dropout strength.
45         - weight_scale: Scalar giving the standard deviation for random
46           initialization of the weights.
47         - reg: Scalar giving L2 regularization strength.
48         """
49         self.params = {}
50         self.reg = reg
51
52         # ===== #
53         # YOUR CODE HERE:
54         # Initialize W1, W2, b1, and b2. Store these as self.params['W1'],
55         # self.params['W2'], self.params['b1'] and self.params['b2']. The
56         # biases are initialized to zero and the weights are initialized
57         # so that each parameter has mean 0 and standard deviation
58         weight_scale.

```

```

57     # The dimensions of W1 should be (input_dim, hidden_dim) and the
58     # dimensions of W2 should be (hidden_dims, num_classes)
59     # ===== #
60
61
62     # ===== #
63     # END YOUR CODE HERE
64     # ===== #
65
66 def loss(self, X, y=None):
67     """
68     Compute loss and gradient for a minibatch of data.
69
70     Inputs:
71     - X: Array of input data of shape (N, d_1, ..., d_k)
72     - y: Array of labels, of shape (N,). y[i] gives the label for X[i].
73
74     Returns:
75     If y is None, then run a test-time forward pass of the model and
return:
76     - scores: Array of shape (N, C) giving classification scores, where
77       scores[i, c] is the classification score for X[i] and class c.
78
79     If y is not None, then run a training-time forward and backward pass
and
80     return a tuple of:
81     - loss: Scalar value giving the loss
82     - grads: Dictionary with the same keys as self.params, mapping
parameter
83       names to gradients of the loss with respect to those parameters.
84     """
85     scores = None
86
87     # ===== #
88     # YOUR CODE HERE:
89     # Implement the forward pass of the two-layer neural network. Store
90     # the class scores as the variable 'scores'. Be sure to use the
layers
91     # you prior implemented.
92     # ===== #
93
94     # ===== #
95     # END YOUR CODE HERE
96     # ===== #
97
98     # If y is None then we are in test mode so just return scores
99     if y is None:
100         return scores
101
102     loss, grads = 0, {}
103     # ===== #
104     # YOUR CODE HERE:
105     # Implement the backward pass of the two-layer neural net. Store
106     # the loss as the variable 'loss' and store the gradients in the
107     # 'grads' dictionary. For the grads dictionary, grads['W1'] holds
108     # the gradient for W1, grads['b1'] holds the gradient for b1, etc.
109     # i.e., grads[k] holds the gradient for self.params[k].
110     #

```

```

111         # Add L2 regularization, where there is an added cost
112         0.5*self.reg*W^2
113         # for each W. Be sure to include the 0.5 multiplying factor to
114         # match our implementation.
115         # And be sure to use the layers you prior implemented.
116         # ===== #
117
118         # ===== #
119         # END YOUR CODE HERE
120         # ===== #
121
122         return loss, grads
123
124
125 class FullyConnectedNet(object):
126     """
127     A fully-connected neural network with an arbitrary number of hidden
128     layers,
129     ReLU nonlinearities, and a softmax loss function. This will also
130     implement
131     dropout and batch normalization as options. For a network with L layers,
132     the architecture will be
133     {affine - [batch norm] - relu - [dropout]} x (L - 1) - affine - softmax
134     where batch normalization and dropout are optional, and the {...} block
135     is
136     repeated L - 1 times.
137
138     Similar to the TwoLayerNet above, learnable parameters are stored in the
139     self.params dictionary and will be learned using the Solver class.
140     """
141     def __init__(self, hidden_dims, input_dim=3*32*32, num_classes=10,
142                 dropout=0, use_batchnorm=False, reg=0.0,
143                 weight_scale=1e-2, dtype=np.float32, seed=None):
144         """
145         Initialize a new FullyConnectedNet.
146
147         Inputs:
148         - hidden_dims: A list of integers giving the size of each hidden
149         layer.
150         - input_dim: An integer giving the size of the input.
151         - num_classes: An integer giving the number of classes to classify.
152         - dropout: Scalar between 0 and 1 giving dropout strength. If
153         dropout=1 then
154         the network should not use dropout at all.
155         - use_batchnorm: Whether or not the network should use batch
156         normalization.
157         - reg: Scalar giving L2 regularization strength.
158         - weight_scale: Scalar giving the standard deviation for random
159         initialization of the weights.
160         - dtype: A numpy datatype object; all computations will be performed
161         using
162         this datatype. float32 is faster but less accurate, so you should
163         use
164         float64 for numeric gradient checking.

```

```

160         - seed: If not None, then pass this random seed to the dropout
layers. This
161         will make the dropout layers deterministic so we can gradient check
the
162         model.
163         """
164         self.use_batchnorm = use_batchnorm
165         self.use_dropout = dropout < 1
166         self.reg = reg
167         self.num_layers = 1 + len(hidden_dims)
168         self.dtype = dtype
169         self.params = {}
170
171         # ===== #
172         # YOUR CODE HERE:
173         # Initialize all parameters of the network in the self.params
dictionary.
174         # The weights and biases of layer 1 are W1 and b1; and in general
the
175         # weights and biases of layer i are Wi and bi. The
176         # biases are initialized to zero and the weights are initialized
177         # so that each parameter has mean 0 and standard deviation
weight_scale.
178         #
179         # BATCHNORM: Initialize the gammas of each layer to 1 and the beta
180         # parameters to zero. The gamma and beta parameters for layer 1
should
181         # be self.params['gamma1'] and self.params['beta1']. For layer 2,
they
182         # should be gamma2 and beta2, etc. Only use batchnorm if
self.use_batchnorm
183         # is true and DO NOT do batch normalize the output scores.
184         # ===== #
185
186         dims = hidden_dims[:]
187         dims.insert(0, input_dim)
188         dims.append(num_classes)
189
190         for i in range(len(dims) - 1):
191             self.params['W{}'.format(i+1)] = weight_scale *
np.random.randn(dims[i], dims[i+1])
192             self.params['b{}'.format(i+1)] = np.zeros(dims[i+1])
193
194         if use_batchnorm:
195             for i in range(self.num_layers - 1):
196                 self.params['gamma{}'.format(i+1)] = np.ones(hidden_dims[i])
197                 self.params['beta{}'.format(i+1)] = np.zeros(hidden_dims[i])
198
199
200         # ===== #
201         # END YOUR CODE HERE
202         # ===== #
203
204         # When using dropout we need to pass a dropout_param dictionary to
each
205         # dropout layer so that the layer knows the dropout probability and
the mode
206         # (train / test). You can pass the same dropout_param to each dropout
layer.
207         self.dropout_param = {}

```

```

208         if self.use_dropout:
209             self.dropout_param = {'mode': 'train', 'p': dropout}
210         if seed is not None:
211             self.dropout_param['seed'] = seed
212
213         # With batch normalization we need to keep track of running means and
214         # variances, so we need to pass a special bn_param object to each
batch
215         # normalization layer. You should pass self.bn_params[0] to the
forward pass
216         # of the first batch normalization layer, self.bn_params[1] to the
forward
217         # pass of the second batch normalization layer, etc.
218         self.bn_params = []
219         if self.use_batchnorm:
220             self.bn_params = [{'mode': 'train'} for i in
np.arange(self.num_layers - 1)]
221
222         # Cast all parameters to the correct datatype
223         for k, v in self.params.items():
224             self.params[k] = v.astype(dtype)
225
226
227     def loss(self, X, y=None):
228         """
229         Compute loss and gradient for the fully-connected net.
230
231         Input / output: Same as TwoLayerNet above.
232         """
233         X = X.astype(self.dtype)
234         mode = 'test' if y is None else 'train'
235
236         # Set train/test mode for batchnorm params and dropout param since
they
237         # behave differently during training and testing.
238         if self.dropout_param is not None:
239             self.dropout_param['mode'] = mode
240         if self.use_batchnorm:
241             for bn_param in self.bn_params:
242                 bn_param[mode] = mode
243
244         scores = None
245
246         # ===== #
247         # YOUR CODE HERE:
248         # Implement the forward pass of the FC net and store the output
249         # scores as the variable "scores".
250         #
251         # BATCHNORM: If self.use_batchnorm is true, insert a bathnorm layer
252         # between the affine_forward and relu_forward layers. You may
253         # also write an affine_batchnorm_relu() function in layer_utils.py.
254         #
255         # DROPOUT: If dropout is non-zero, insert a dropout layer after
256         # every ReLU layer.
257         # ===== #
258
259         caches = []
260         dropout_caches = []
261         x = X
262         for i in range(1, self.num_layers + 1):

```

```

263         w, b = self.params['W{}'.format(i)], self.params['b{}'.format(i)]
264         if (i == self.num_layers):
265             scores, cache = affine_forward(x, w, b)
266         else:
267             if self.use_batchnorm:
268                 x, cache = affine_batchnorm_relu_forward(x, w, b,
self.params['gamma{}'.format(i)], self.params['beta{}'.format(i)],
self.bn_params[i-1])
269             else:
270                 x, cache = affine_relu_forward(x, w, b)
271
272             if self.use_dropout:
273                 x, dropout_cache = dropout_forward(x, self.dropout_param)
274                 dropout_caches.append(dropout_cache)
275             caches.append(cache)
276
277             # ===== #
278             # END YOUR CODE HERE
279             # ===== #
280
281             # If test mode return early
282             if mode == 'test':
283                 return scores
284
285             loss, grads = 0.0, {}
286             # ===== #
287             # YOUR CODE HERE:
288             # Implement the backwards pass of the FC net and store the
gradients
289             # in the grads dict, so that grads[k] is the gradient of
self.params[k]
290             # Be sure your L2 regularization includes a 0.5 factor.
291             #
292             # BATCHNORM: Incorporate the backward pass of the batchnorm.
293             #
294             # DROPOUT: Incorporate the backward pass of dropout.
295             # ===== #
296
297
298             sm_loss, dout = softmax_loss(scores, y)
299             reg_loss = 0.0
300             for i in range(1, self.num_layers + 1):
301                 reg_loss += 0.5 * self.reg *
(np.sum(self.params['W{}'.format(i)]**2))
302             loss = sm_loss + reg_loss
303
304             for i in range(self.num_layers, 0, -1):
305                 cur_w = 'W{}'.format(i)
306                 cur_b = 'b{}'.format(i)
307                 if (i == self.num_layers):
308                     dout, grads[cur_w], grads[cur_b] = affine_backward(dout,
caches.pop())
309                 else:
310                     if self.use_dropout:
311                         dout = dropout_backward(dout, dropout_caches.pop())
312                     if self.use_batchnorm:
313                         dout, grads[cur_w], grads[cur_b], grads['gamma{}'.format(i)],
grads['beta{}'.format(i)] = affine_batchnorm_relu_backward(dout,
caches.pop())
314                 else:

```

```
315         dout, grads[cur_w], grads[cur_b] = affine_relu_backward(dout,
316 caches.pop())
317         grads[cur_w] += self.reg * self.params[cur_w]
318         # ===== #
319         # END YOUR CODE HERE
320         # ===== #
321
322     return loss, grads
323
```