

CNN-BatchNorm

February 24, 2020

0.1 Spatial batch normalization

In fully connected networks, we performed batch normalization on the activations. To do something equivalent on CNNs, we modify batch normalization slightly.

Normally batch-normalization accepts inputs of shape (N, D) and produces outputs of shape (N, D) , where we normalize across the minibatch dimension N . For data coming from convolutional layers, batch normalization accepts inputs of shape (N, C, H, W) and produces outputs of shape (N, C, H, W) where the N dimension gives the minibatch size and the (H, W) dimensions give the spatial size of the feature map.

How do we calculate the spatial averages? First, notice that for the C feature maps we have (i.e., the layer has C filters) that each of these ought to have its own batch norm statistics, since each feature map may be picking out very different features in the images. However, within a feature map, we may assume that across all inputs and across all locations in the feature map, there ought to be relatively similar first and second order statistics. Hence, one way to think of spatial batch-normalization is to reshape the (N, C, H, W) array as an $(N*H*W, C)$ array and perform batch normalization on this array.

Since spatial batch norm and batch normalization are similar, it'd be good to at this point also copy and paste our prior implemented layers from HW #4. Please copy and paste your prior implemented code from HW #4 to start this assignment. If you did not correctly implement the layers in HW #4, you may collaborate with a classmate to use their implementations from HW #4. You may also visit TA or Prof OH to correct your implementation.

You'll want to copy and paste from HW #4: - layers.py for your FC network layers, as well as batchnorm and dropout. - layer_utils.py for your combined FC network layers. - optim.py for your optimizers.

Be sure to place these in the nndl/ directory so they're imported correctly. Note, as announced in class, we will not be releasing our solutions.

If you use your prior implementations of the batchnorm, then your spatial batchnorm implementation may be very short. Our implementations of the forward and backward pass are each 6 lines of code.

CS231n has built a solid API for building these modular frameworks and training them, and we will use their very well implemented framework as opposed to "reinventing the wheel." This includes using their Solver, various utility functions, their layer structure, and their implementation of fast CNN layers. This also includes nndl.fc_net, nndl.layers, and nndl.layer_utils. As in prior assignments, we thank Serena Yeung & Justin Johnson for permission to use code written for the CS 231n class (cs231n.stanford.edu).

```
In [1]: ## Import and setups
```

```

import time
import numpy as np
import matplotlib.pyplot as plt
from nndl.conv_layers import *
from cs231n.data_utils import get_CIFAR10_data
from cs231n.gradient_check import eval_numerical_gradient, eval_numerical_gradient_array
from cs231n.solver import Solver

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2

def rel_error(x, y):
    """ returns relative error """
    return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))

```

0.2 Spatial batch normalization forward pass

Implement the forward pass, `spatial_batchnorm_forward` in `nndl/conv_layers.py`. Test your implementation by running the cell below.

In [2]: *# Check the training-time forward pass by checking means and variances of features both before and after spatial batch normalization*

```

N, C, H, W = 2, 3, 4, 5
x = 4 * np.random.randn(N, C, H, W) + 10

print('Before spatial batch normalization:')
print('  Shape: ', x.shape)
print('  Means: ', x.mean(axis=(0, 2, 3)))
print('  Stds: ', x.std(axis=(0, 2, 3)))

# Means should be close to zero and stds close to one
gamma, beta = np.ones(C), np.zeros(C)
bn_param = {'mode': 'train'}
out, _ = spatial_batchnorm_forward(x, gamma, beta, bn_param)
print('After spatial batch normalization:')
print('  Shape: ', out.shape)
print('  Means: ', out.mean(axis=(0, 2, 3)))
print('  Stds: ', out.std(axis=(0, 2, 3)))

# Means should be close to beta and stds close to gamma

```

```

gamma, beta = np.asarray([3, 4, 5]), np.asarray([6, 7, 8])
out, _ = spatial_batchnorm_forward(x, gamma, beta, bn_param)
print('After spatial batch normalization (nontrivial gamma, beta):')
print('  Shape: ', out.shape)
print('  Means: ', out.mean(axis=(0, 2, 3)))
print('  Stds: ', out.std(axis=(0, 2, 3)))

```

Before spatial batch normalization:

```

Shape: (2, 3, 4, 5)
Means: [10.31867256 11.24502102  9.0823411 ]
Stds:  [3.52710979 3.91035495 4.15873526]

```

After spatial batch normalization:

```

Shape: (2, 3, 4, 5)
Means: [ 4.39925874e-16  1.15879528e-16 -4.74620343e-16]
Stds:  [0.99999996 0.99999967 0.99999971]

```

After spatial batch normalization (nontrivial gamma, beta):

```

Shape: (2, 3, 4, 5)
Means: [6. 7. 8.]
Stds:  [2.99999879 3.99999869 4.99999855]

```

0.3 Spatial batch normalization backward pass

Implement the backward pass, `spatial_batchnorm_backward` in `nndl/conv_layers.py`. Test your implementation by running the cell below.

```

In [3]: N, C, H, W = 2, 3, 4, 5
        x = 5 * np.random.randn(N, C, H, W) + 12
        gamma = np.random.randn(C)
        beta = np.random.randn(C)
        dout = np.random.randn(N, C, H, W)

        bn_param = {'mode': 'train'}
        fx = lambda x: spatial_batchnorm_forward(x, gamma, beta, bn_param)[0]
        fg = lambda a: spatial_batchnorm_forward(x, gamma, beta, bn_param)[0]
        fb = lambda b: spatial_batchnorm_forward(x, gamma, beta, bn_param)[0]

        dx_num = eval_numerical_gradient_array(fx, x, dout)
        da_num = eval_numerical_gradient_array(fg, gamma, dout)
        db_num = eval_numerical_gradient_array(fb, beta, dout)

        _, cache = spatial_batchnorm_forward(x, gamma, beta, bn_param)
        dx, dgamma, dbeta = spatial_batchnorm_backward(dout, cache)
        print('dx error: ', rel_error(dx_num, dx))
        print('dgamma error: ', rel_error(da_num, dgamma))
        print('dbeta error: ', rel_error(db_num, dbeta))

dx error: 3.6730049732661024e-08
dgamma error: 7.721967973774263e-12

```

dbeta error: 3.275594161332795e-12

CNN-Layers

February 24, 2020

0.1 Convolutional neural network layers

In this notebook, we will build the convolutional neural network layers. This will be followed by a spatial batchnorm, and then in the final notebook of this assignment, we will train a CNN to further improve the validation accuracy on CIFAR-10.

CS231n has built a solid API for building these modular frameworks and training them, and we will use their very well implemented framework as opposed to “reinventing the wheel.” This includes using their Solver, various utility functions, their layer structure, and their implementation of fast CNN layers. This also includes `nndl.fc_net`, `nndl.layers`, and `nndl.layer_utils`. As in prior assignments, we thank Serena Yeung & Justin Johnson for permission to use code written for the CS 231n class (cs231n.stanford.edu).

In [1]: *## Import and setups*

```
import time
import numpy as np
import matplotlib.pyplot as plt
from nndl.conv_layers import *
from cs231n.data_utils import get_CIFAR10_data
from cs231n.gradient_check import eval_numerical_gradient, eval_numerical_gradient_array
from cs231n.solver import Solver

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2

def rel_error(x, y):
    """ returns relative error """
    return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))
```

0.2 Implementing CNN layers

Just as we implemented modular layers for fully connected networks, batch normalization, and dropout, we'll want to implement modular layers for convolutional neural networks. These layers are in `nndl/conv_layers.py`.

0.2.1 Convolutional forward pass

Begin by implementing a naive version of the forward pass of the CNN that uses for loops. This function is `conv_forward_naive` in `nndl/conv_layers.py`. Don't worry about efficiency of implementation. Later on, we provide a fast implementation of these layers. This version ought to test your understanding of convolution. In our implementation, there is a triple for loop.

After you implement `conv_forward_naive`, test your implementation by running the cell below.

```
In [2]: x_shape = (2, 3, 4, 4)
        w_shape = (3, 3, 4, 4)
        x = np.linspace(-0.1, 0.5, num=np.prod(x_shape)).reshape(x_shape)
        w = np.linspace(-0.2, 0.3, num=np.prod(w_shape)).reshape(w_shape)
        b = np.linspace(-0.1, 0.2, num=3)

        conv_param = {'stride': 2, 'pad': 1}
        out, _ = conv_forward_naive(x, w, b, conv_param)
        correct_out = np.array([[[[-0.08759809, -0.10987781],
                                   [-0.18387192, -0.2109216 ]],
                                [[ 0.21027089,  0.21661097],
                                   [ 0.22847626,  0.23004637]],
                                [[ 0.50813986,  0.54309974],
                                   [ 0.64082444,  0.67101435]]],
                               [[[-0.98053589, -1.03143541],
                                   [-1.19128892, -1.24695841]],
                                [[ 0.69108355,  0.66880383],
                                   [ 0.59480972,  0.56776003]],
                                [[ 2.36270298,  2.36904306],
                                   [ 2.38090835,  2.38247847]]]])

        # Compare your output to ours; difference should be around 1e-8
        print('Testing conv_forward_naive')
        print('difference: ', rel_error(out, correct_out))
```

```
Testing conv_forward_naive
difference:  2.2121476417505994e-08
```

0.2.2 Convolutional backward pass

Now, implement a naive version of the backward pass of the CNN. The function is `conv_backward_naive` in `nndl/conv_layers.py`. Don't worry about efficiency of implementation. Later on, we provide a fast implementation of these layers. This version ought to test your understanding of convolution. In our implementation, there is a quadruple for loop.

After you implement `conv_backward_naive`, test your implementation by running the cell below.

```
In [3]: x = np.random.randn(4, 3, 5, 5)
        w = np.random.randn(2, 3, 3, 3)
        b = np.random.randn(2,)
        dout = np.random.randn(4, 2, 5, 5)
        conv_param = {'stride': 1, 'pad': 1}

        out, cache = conv_forward_naive(x,w,b,conv_param)

        dx_num = eval_numerical_gradient_array(lambda x: conv_forward_naive(x, w, b, conv_param), x, dout)
        dw_num = eval_numerical_gradient_array(lambda w: conv_forward_naive(x, w, b, conv_param), w, dout)
        db_num = eval_numerical_gradient_array(lambda b: conv_forward_naive(x, w, b, conv_param), b, dout)

        out, cache = conv_forward_naive(x, w, b, conv_param)
        dx, dw, db = conv_backward_naive(dout, cache)

        # Your errors should be around 1e-9'
        print('Testing conv_backward_naive function')
        print('dx error: ', rel_error(dx, dx_num))
        print('dw error: ', rel_error(dw, dw_num))
        print('db error: ', rel_error(db, db_num))
```

```
Testing conv_backward_naive function
dx error:  1.1013003669713249e-08
dw error:  4.724119625184007e-10
db error:  1.1266267530410223e-11
```

0.2.3 Max pool forward pass

In this section, we will implement the forward pass of the max pool. The function is `max_pool_forward_naive` in `nndl/conv_layers.py`. Do not worry about the efficiency of implementation.

After you implement `max_pool_forward_naive`, test your implementation by running the cell below.

```
In [4]: x_shape = (2, 3, 4, 4)
        x = np.linspace(-0.3, 0.4, num=np.prod(x_shape)).reshape(x_shape)
        pool_param = {'pool_width': 2, 'pool_height': 2, 'stride': 2}

        out, _ = max_pool_forward_naive(x, pool_param)

        correct_out = np.array([[[[-0.26315789, -0.24842105],
                                   [-0.20421053, -0.18947368]],
                                  [[-0.14526316, -0.13052632],
                                   [-0.08631579, -0.07157895]],
                                  [[-0.02736842, -0.01263158],
```

```

[ 0.03157895, 0.04631579]]],
[[[ 0.09052632, 0.10526316],
[ 0.14947368, 0.16421053]],
[[ 0.20842105, 0.22315789],
[ 0.26736842, 0.28210526]],
[[ 0.32631579, 0.34105263],
[ 0.38526316, 0.4         ]]]])

```

```

# Compare your output with ours. Difference should be around 1e-8.
print('Testing max_pool_forward_naive function:')
print('difference: ', rel_error(out, correct_out))

```

```

Testing max_pool_forward_naive function:
difference: 4.1666665157267834e-08

```

0.2.4 Max pool backward pass

In this section, you will implement the backward pass of the max pool. The function is `max_pool_backward_naive` in `nndl/conv_layers.py`. Do not worry about the efficiency of implementation.

After you implement `max_pool_backward_naive`, test your implementation by running the cell below.

```

In [5]: x = np.random.randn(3, 2, 8, 8)
        dout = np.random.randn(3, 2, 4, 4)
        pool_param = {'pool_height': 2, 'pool_width': 2, 'stride': 2}

        dx_num = eval_numerical_gradient_array(lambda x: max_pool_forward_naive(x, pool_param)

        out, cache = max_pool_forward_naive(x, pool_param)
        dx = max_pool_backward_naive(dout, cache)

        # Your error should be around 1e-12
        print('Testing max_pool_backward_naive function:')
        print('dx error: ', rel_error(dx, dx_num))

```

```

Testing max_pool_backward_naive function:
dx error: 3.275629170287585e-12

```

0.3 Fast implementation of the CNN layers

Implementing fast versions of the CNN layers can be difficult. We will provide you with the fast layers implemented by `cs231n`. They are provided in `cs231n/fast_layers.py`.

The fast convolution implementation depends on a Cython extension; to compile it you need to run the following from the `cs231n` directory:

```
python setup.py build_ext --inplace
```


NOTE: The fast implementation for pooling will only perform optimally if the pooling regions are non-overlapping and tile the input. If these conditions are not met then the fast pooling implementation will not be much faster than the naive implementation.

You can compare the performance of the naive and fast versions of these layers by running the cell below.

You should see pretty drastic speedups in the implementation of these layers. On our machine, the forward pass speeds up by 17x and the backward pass speeds up by 840x. Of course, these numbers will vary from machine to machine, as well as on your precise implementation of the naive layers.

```
In [6]: from cs231n.fast_layers import conv_forward_fast, conv_backward_fast
        from time import time
```

```
x = np.random.randn(100, 3, 31, 31)
w = np.random.randn(25, 3, 3, 3)
b = np.random.randn(25,)
dout = np.random.randn(100, 25, 16, 16)
conv_param = {'stride': 2, 'pad': 1}

t0 = time()
out_naive, cache_naive = conv_forward_naive(x, w, b, conv_param)
t1 = time()
out_fast, cache_fast = conv_forward_fast(x, w, b, conv_param)
t2 = time()

print('Testing conv_forward_fast:')
print('Naive: %fs' % (t1 - t0))
print('Fast: %fs' % (t2 - t1))
print('Speedup: %fx' % ((t1 - t0) / (t2 - t1)))
print('Difference: ', rel_error(out_naive, out_fast))

t0 = time()
dx_naive, dw_naive, db_naive = conv_backward_naive(dout, cache_naive)
t1 = time()
dx_fast, dw_fast, db_fast = conv_backward_fast(dout, cache_fast)
t2 = time()

print('\nTesting conv_backward_fast:')
print('Naive: %fs' % (t1 - t0))
print('Fast: %fs' % (t2 - t1))
print('Speedup: %fx' % ((t1 - t0) / (t2 - t1)))
print('dx difference: ', rel_error(dx_naive, dx_fast))
print('dw difference: ', rel_error(dw_naive, dw_fast))
print('db difference: ', rel_error(db_naive, db_fast))
```

Testing conv_forward_fast:

Naive: 10.333672s

Fast: 0.047794s

Speedup: 216.211195x
Difference: 4.649974808364135e-12

Testing conv_backward_fast:
Naive: 12.039610s
Fast: 0.016100s
Speedup: 747.805118x
dx difference: 7.823110457071484e-12
dw difference: 4.491393368646801e-13
db difference: 0.0

```
In [7]: from cs231n.fast_layers import max_pool_forward_fast, max_pool_backward_fast
```

```
x = np.random.randn(100, 3, 32, 32)
dout = np.random.randn(100, 3, 16, 16)
pool_param = {'pool_height': 2, 'pool_width': 2, 'stride': 2}

t0 = time()
out_naive, cache_naive = max_pool_forward_naive(x, pool_param)
t1 = time()
out_fast, cache_fast = max_pool_forward_fast(x, pool_param)
t2 = time()

print('Testing pool_forward_fast:')
print('Naive: %fs' % (t1 - t0))
print('fast: %fs' % (t2 - t1))
print('speedup: %fx' % ((t1 - t0) / (t2 - t1)))
print('difference: ', rel_error(out_naive, out_fast))

t0 = time()
dx_naive = max_pool_backward_naive(dout, cache_naive)
t1 = time()
dx_fast = max_pool_backward_fast(dout, cache_fast)
t2 = time()

print('\nTesting pool_backward_fast:')
print('Naive: %fs' % (t1 - t0))
print('speedup: %fx' % ((t1 - t0) / (t2 - t1)))
print('dx difference: ', rel_error(dx_naive, dx_fast))
```

Testing pool_forward_fast:
Naive: 0.694415s
fast: 0.005271s
speedup: 131.743577x
difference: 0.0

Testing pool_backward_fast:

```
Naive: 1.862069s
speedup: 123.940078x
dx difference: 0.0
```

0.4 Implementation of cascaded layers

We've provided the following functions in `nndl/conv_layer_utils.py`: - `conv_relu_forward` - `conv_relu_backward` - `conv_relu_pool_forward` - `conv_relu_pool_backward`

These use the fast implementations of the conv net layers. You can test them below:

```
In [8]: from nndl.conv_layer_utils import conv_relu_pool_forward, conv_relu_pool_backward
```

```
x = np.random.randn(2, 3, 16, 16)
w = np.random.randn(3, 3, 3, 3)
b = np.random.randn(3,)
dout = np.random.randn(2, 3, 8, 8)
conv_param = {'stride': 1, 'pad': 1}
pool_param = {'pool_height': 2, 'pool_width': 2, 'stride': 2}
```

```
out, cache = conv_relu_pool_forward(x, w, b, conv_param, pool_param)
dx, dw, db = conv_relu_pool_backward(dout, cache)
```

```
dx_num = eval_numerical_gradient_array(lambda x: conv_relu_pool_forward(x, w, b, conv_param, pool_param), x, dout)
dw_num = eval_numerical_gradient_array(lambda w: conv_relu_pool_forward(x, w, b, conv_param, pool_param), w, dout)
db_num = eval_numerical_gradient_array(lambda b: conv_relu_pool_forward(x, w, b, conv_param, pool_param), b, dout)
```

```
print('Testing conv_relu_pool')
print('dx error: ', rel_error(dx_num, dx))
print('dw error: ', rel_error(dw_num, dw))
print('db error: ', rel_error(db_num, db))
```

```
Testing conv_relu_pool
dx error: 2.0047909130494352e-08
dw error: 3.710881095531974e-09
db error: 3.1888427581489564e-10
```

```
In [9]: from nndl.conv_layer_utils import conv_relu_forward, conv_relu_backward
```

```
x = np.random.randn(2, 3, 8, 8)
w = np.random.randn(3, 3, 3, 3)
b = np.random.randn(3,)
dout = np.random.randn(2, 3, 8, 8)
conv_param = {'stride': 1, 'pad': 1}
```

```
out, cache = conv_relu_forward(x, w, b, conv_param)
dx, dw, db = conv_relu_backward(dout, cache)
```

```

dx_num = eval_numerical_gradient_array(lambda x: conv_relu_forward(x, w, b, conv_param),
dw_num = eval_numerical_gradient_array(lambda w: conv_relu_forward(x, w, b, conv_param),
db_num = eval_numerical_gradient_array(lambda b: conv_relu_forward(x, w, b, conv_param),

print('Testing conv_relu:')
print('dx error: ', rel_error(dx_num, dx))
print('dw error: ', rel_error(dw_num, dw))
print('db error: ', rel_error(db_num, db))

```

```

Testing conv_relu:
dx error:  9.191014012570723e-10
dw error:  1.7710199249080962e-09
db error:  9.667399468598408e-12

```

0.5 What next?

We saw how helpful batch normalization was for training FC nets. In the next notebook, we'll implement a batch normalization for convolutional neural networks, and then finish off by implementing a CNN to improve our validation accuracy on CIFAR-10.

CNN

February 25, 2020

1 Convolutional neural networks

In this notebook, we'll put together our convolutional layers to implement a 3-layer CNN. Then, we'll ask you to implement a CNN that can achieve $> 65\%$ validation error on CIFAR-10.

CS231n has built a solid API for building these modular frameworks and training them, and we will use their very well implemented framework as opposed to “reinventing the wheel.” This includes using their Solver, various utility functions, their layer structure, and their implementation of fast CNN layers. This also includes `nndl.fc_net`, `nndl.layers`, and `nndl.layer_utils`. As in prior assignments, we thank Serena Yeung & Justin Johnson for permission to use code written for the CS 231n class (cs231n.stanford.edu).

If you have not completed the Spatial BatchNorm Notebook, please see the following description from that notebook:

Please copy and paste your prior implemented code from HW #4 to start this assignment. If you did not correctly implement the layers in HW #4, you may collaborate with a classmate to use their layer implementations from HW #4. You may also visit TA or Prof OH to correct your implementation.

You'll want to copy and paste from HW #4: - `layers.py` for your FC network layers, as well as `batchnorm` and `dropout`. - `layer_utils.py` for your combined FC network layers. - `optim.py` for your optimizers.

Be sure to place these in the `nndl/` directory so they're imported correctly. Note, as announced in class, we will not be releasing our solutions.

```
In [1]: # As usual, a bit of setup
```

```
import numpy as np
import matplotlib.pyplot as plt
from nndl.cnn import *
from cs231n.data_utils import get_CIFAR10_data
from cs231n.gradient_check import eval_numerical_gradient_array, eval_numerical_gradients
from nndl.layers import *
from nndl.conv_layers import *
from cs231n.fast_layers import *
from cs231n.solver import Solver

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
```

```

plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2

def rel_error(x, y):
    """ returns relative error """
    return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))

In [2]: # Load the (preprocessed) CIFAR10 data.

data = get_CIFAR10_data()
for k in data.keys():
    print('{}: {}'.format(k, data[k].shape))

X_train: (49000, 3, 32, 32)
y_train: (49000,)
X_val: (1000, 3, 32, 32)
y_val: (1000,)
X_test: (1000, 3, 32, 32)
y_test: (1000,)

```

1.1 Three layer CNN

In this notebook, you will implement a three layer CNN. The `ThreeLayerConvNet` class is in `mndl/cnn.py`. You'll need to modify that code for this section, including the initialization, as well as the calculation of the loss and gradients. You should be able to use the building blocks you have either earlier coded or that we have provided. Be sure to use the fast layers.

The architecture of this CNN will be:

conv - relu - 2x2 max pool - affine - relu - affine - softmax

We won't use batchnorm yet. You've also done enough of these to know how to debug; use the cells below.

Note: As we are implementing several layers CNN networks. The gradient error can be expected for the `eval_numerical_gradient()` function. If your `W1` max relative error and `W2` max relative error are around or below 0.01, they should be acceptable. Other errors should be less than $1e-5$.

```

In [3]: num_inputs = 2
        input_dim = (3, 16, 16)
        reg = 0.0
        num_classes = 10
        X = np.random.randn(num_inputs, *input_dim)
        y = np.random.randint(num_classes, size=num_inputs)

        model = ThreeLayerConvNet(num_filters=3, filter_size=3,
                                   input_dim=input_dim, hidden_dim=7,

```

```

dtype=np.float64)
loss, grads = model.loss(X, y)
for param_name in sorted(grads):
    f = lambda _: model.loss(X, y)[0]
    param_grad_num = eval_numerical_gradient(f, model.params[param_name], verbose=False)
    e = rel_error(param_grad_num, grads[param_name])
    print('{} max relative error: {}'.format(param_name, rel_error(param_grad_num, grads[param_name])))

W1 max relative error: 0.0006507149025299291
W2 max relative error: 0.00568316345265749
W3 max relative error: 4.2889602827172216e-05
b1 max relative error: 1.7735505300977068e-05
b2 max relative error: 8.581484252488379e-07
b3 max relative error: 1.1508232107915781e-09

```

1.1.1 Overfit small dataset

To check your CNN implementation, let's overfit a small dataset.

```

In [4]: num_train = 100
        small_data = {
            'X_train': data['X_train'][:num_train],
            'y_train': data['y_train'][:num_train],
            'X_val': data['X_val'],
            'y_val': data['y_val'],
        }

        model = ThreeLayerConvNet(weight_scale=1e-2)

        solver = Solver(model, small_data,
                        num_epochs=10, batch_size=50,
                        update_rule='adam',
                        optim_config={
                            'learning_rate': 1e-3,
                        },
                        verbose=True, print_every=1)

        solver.train()


(Iteration 1 / 20) loss: 2.381830
(Epoch 0 / 10) train acc: 0.160000; val_acc: 0.119000
(Iteration 2 / 20) loss: 5.928924
(Epoch 1 / 10) train acc: 0.190000; val_acc: 0.130000
(Iteration 3 / 20) loss: 2.551710
(Iteration 4 / 20) loss: 2.735981
(Epoch 2 / 10) train acc: 0.180000; val_acc: 0.079000
(Iteration 5 / 20) loss: 3.169556
(Iteration 6 / 20) loss: 2.329534
(Epoch 3 / 10) train acc: 0.260000; val_acc: 0.169000

```

```
(Iteration 7 / 20) loss: 2.298394
(Iteration 8 / 20) loss: 1.884549
(Epoch 4 / 10) train acc: 0.380000; val_acc: 0.152000
(Iteration 9 / 20) loss: 1.923281
(Iteration 10 / 20) loss: 1.841728
(Epoch 5 / 10) train acc: 0.470000; val_acc: 0.155000
(Iteration 11 / 20) loss: 1.494036
(Iteration 12 / 20) loss: 1.623882
(Epoch 6 / 10) train acc: 0.590000; val_acc: 0.190000
(Iteration 13 / 20) loss: 1.293669
(Iteration 14 / 20) loss: 1.343754
(Epoch 7 / 10) train acc: 0.620000; val_acc: 0.191000
(Iteration 15 / 20) loss: 1.286494
(Iteration 16 / 20) loss: 0.993775
(Epoch 8 / 10) train acc: 0.630000; val_acc: 0.204000
(Iteration 17 / 20) loss: 0.985476
(Iteration 18 / 20) loss: 0.925707
(Epoch 9 / 10) train acc: 0.680000; val_acc: 0.243000
(Iteration 19 / 20) loss: 1.135033
(Iteration 20 / 20) loss: 0.856049
(Epoch 10 / 10) train acc: 0.800000; val_acc: 0.205000
```

```
In [5]: plt.subplot(2, 1, 1)
        plt.plot(solver.loss_history, 'o')
        plt.xlabel('iteration')
        plt.ylabel('loss')

        plt.subplot(2, 1, 2)
        plt.plot(solver.train_acc_history, '-o')
        plt.plot(solver.val_acc_history, '-o')
        plt.legend(['train', 'val'], loc='upper left')
        plt.xlabel('epoch')
        plt.ylabel('accuracy')
        plt.show()
```

1.2 Train the network

Now we train the 3 layer CNN on CIFAR-10 and assess its accuracy.

```
In [6]: model = ThreeLayerConvNet(weight_scale=0.001, hidden_dim=500, reg=0.001)
```

```

solver = Solver(model, data,
                 num_epochs=1, batch_size=50,
                 update_rule='adam',
                 optim_config={
                     'learning_rate': 1e-3,
                 },
                 verbose=True, print_every=20)
solver.train()

```

```

(Iteration 1 / 980) loss: 2.304654
(Epoch 0 / 1) train acc: 0.102000; val_acc: 0.093000
(Iteration 21 / 980) loss: 2.131816
(Iteration 41 / 980) loss: 2.020851
(Iteration 61 / 980) loss: 2.117923
(Iteration 81 / 980) loss: 2.010330
(Iteration 101 / 980) loss: 1.687490
(Iteration 121 / 980) loss: 1.882670
(Iteration 141 / 980) loss: 1.939549
(Iteration 161 / 980) loss: 1.552810

```

```
(Iteration 181 / 980) loss: 1.823595
(Iteration 201 / 980) loss: 1.808657
(Iteration 221 / 980) loss: 1.854325
(Iteration 241 / 980) loss: 1.761188
(Iteration 261 / 980) loss: 1.800233
(Iteration 281 / 980) loss: 1.567720
(Iteration 301 / 980) loss: 1.539409
(Iteration 321 / 980) loss: 2.138078
(Iteration 341 / 980) loss: 1.746798
(Iteration 361 / 980) loss: 1.840328
(Iteration 381 / 980) loss: 1.881169
(Iteration 401 / 980) loss: 1.591436
(Iteration 421 / 980) loss: 1.656724
(Iteration 441 / 980) loss: 1.778309
(Iteration 461 / 980) loss: 1.720966
(Iteration 481 / 980) loss: 1.856955
(Iteration 501 / 980) loss: 1.624138
(Iteration 521 / 980) loss: 1.587256
(Iteration 541 / 980) loss: 1.542011
(Iteration 561 / 980) loss: 1.555545
(Iteration 581 / 980) loss: 1.396170
(Iteration 601 / 980) loss: 1.293165
(Iteration 621 / 980) loss: 1.270821
(Iteration 641 / 980) loss: 1.691648
(Iteration 661 / 980) loss: 1.570260
(Iteration 681 / 980) loss: 1.743356
(Iteration 701 / 980) loss: 1.474436
(Iteration 721 / 980) loss: 1.414279
(Iteration 741 / 980) loss: 1.339923
(Iteration 761 / 980) loss: 1.782336
(Iteration 781 / 980) loss: 1.740092
(Iteration 801 / 980) loss: 1.427833
(Iteration 821 / 980) loss: 1.437471
(Iteration 841 / 980) loss: 1.542376
(Iteration 861 / 980) loss: 1.576817
(Iteration 881 / 980) loss: 1.577748
(Iteration 901 / 980) loss: 1.722746
(Iteration 921 / 980) loss: 1.507274
(Iteration 941 / 980) loss: 1.371046
(Iteration 961 / 980) loss: 1.376196
(Epoch 1 / 1) train acc: 0.428000; val_acc: 0.427000
```

2 Get > 65% validation accuracy on CIFAR-10.

In the last part of the assignment, we'll now ask you to train a CNN to get better than 65% validation accuracy on CIFAR-10.

2.0.1 Things you should try:

- Filter size: Above we used 7x7; but VGGNet and onwards showed stacks of 3x3 filters are good.
- Number of filters: Above we used 32 filters. Do more or fewer do better?
- Batch normalization: Try adding spatial batch normalization after convolution layers and vanilla batch normalization after affine layers. Do your networks train faster?
- Network architecture: Can a deeper CNN do better? Consider these architectures:
 - [conv-relu-pool]xN - conv - relu - [affine]xM - [softmax or SVM]
 - [conv-relu-pool]xN - [affine]xM - [softmax or SVM]
 - [conv-relu-conv-relu-pool]xN - [affine]xM - [softmax or SVM]

2.0.2 Tips for training

For each network architecture that you try, you should tune the learning rate and regularization strength. When doing this there are a couple of important things to keep in mind:

- If the parameters are working well, you should see improvement within a few hundred iterations
- Remember the coarse-to-fine approach for hyperparameter tuning: start by testing a large range of hyperparameters for just a few training iterations to find the combinations of parameters that are working at all.
- Once you have found some sets of parameters that seem to work, search more finely around these parameters. You may need to train for more epochs.

```
In [9]: # ===== #
# YOUR CODE HERE:
#   Implement a CNN to achieve greater than 65% validation accuracy
#   on CIFAR-10.
# ===== #

optimizer = 'adam'

learning_rate = 1e-3
lr_decay = 0.9
num_filters = 64
filter_size = 7

model = ThreeLayerConvNet(weight_scale=0.001, hidden_dim=500, reg=0.001)

solver = Solver(model, data,
                 num_epochs=20, batch_size=100,
                 update_rule=optimizer,
                 optim_config={
                     'learning_rate': learning_rate,
                 },
                 lr_decay=lr_decay,
```

```

        verbose=True, print_every=50)
solver.train()

# ===== #
# END YOUR CODE HERE
# ===== #

(Iteration 1 / 9800) loss: 2.304452
(Epoch 0 / 20) train acc: 0.116000; val_acc: 0.105000
(Iteration 51 / 9800) loss: 1.853642
(Iteration 101 / 9800) loss: 1.754378
(Iteration 151 / 9800) loss: 1.585016
(Iteration 201 / 9800) loss: 1.665315
(Iteration 251 / 9800) loss: 1.494164
(Iteration 301 / 9800) loss: 1.452221
(Iteration 351 / 9800) loss: 1.353152
(Iteration 401 / 9800) loss: 1.504349
(Iteration 451 / 9800) loss: 1.365023
(Epoch 1 / 20) train acc: 0.466000; val_acc: 0.491000
(Iteration 501 / 9800) loss: 1.385573
(Iteration 551 / 9800) loss: 1.383519
(Iteration 601 / 9800) loss: 1.342313
(Iteration 651 / 9800) loss: 1.418668
(Iteration 701 / 9800) loss: 1.202109
(Iteration 751 / 9800) loss: 1.360571
(Iteration 801 / 9800) loss: 1.141037
(Iteration 851 / 9800) loss: 1.203072
(Iteration 901 / 9800) loss: 1.451928
(Iteration 951 / 9800) loss: 1.349903
(Epoch 2 / 20) train acc: 0.530000; val_acc: 0.550000
(Iteration 1001 / 9800) loss: 1.270969
(Iteration 1051 / 9800) loss: 1.315878
(Iteration 1101 / 9800) loss: 1.212279
(Iteration 1151 / 9800) loss: 1.474329
(Iteration 1201 / 9800) loss: 1.144212
(Iteration 1251 / 9800) loss: 1.383498
(Iteration 1301 / 9800) loss: 1.388701
(Iteration 1351 / 9800) loss: 1.153182
(Iteration 1401 / 9800) loss: 1.103188
(Iteration 1451 / 9800) loss: 1.214557
(Epoch 3 / 20) train acc: 0.580000; val_acc: 0.578000
(Iteration 1501 / 9800) loss: 1.166137
(Iteration 1551 / 9800) loss: 1.163752
(Iteration 1601 / 9800) loss: 1.019426
(Iteration 1651 / 9800) loss: 1.259598
(Iteration 1701 / 9800) loss: 1.119750
(Iteration 1751 / 9800) loss: 1.221735
(Iteration 1801 / 9800) loss: 1.102449

```

(Iteration 1851 / 9800) loss: 1.122105
(Iteration 1901 / 9800) loss: 1.149631
(Iteration 1951 / 9800) loss: 1.077087
(Epoch 4 / 20) train acc: 0.650000; val_acc: 0.619000
(Iteration 2001 / 9800) loss: 0.993838
(Iteration 2051 / 9800) loss: 1.002742
(Iteration 2101 / 9800) loss: 0.910465
(Iteration 2151 / 9800) loss: 1.000292
(Iteration 2201 / 9800) loss: 0.944183
(Iteration 2251 / 9800) loss: 1.031122
(Iteration 2301 / 9800) loss: 0.908113
(Iteration 2351 / 9800) loss: 1.062253
(Iteration 2401 / 9800) loss: 1.203044
(Epoch 5 / 20) train acc: 0.664000; val_acc: 0.616000
(Iteration 2451 / 9800) loss: 1.003779
(Iteration 2501 / 9800) loss: 1.040453
(Iteration 2551 / 9800) loss: 0.933268
(Iteration 2601 / 9800) loss: 0.988178
(Iteration 2651 / 9800) loss: 1.014140
(Iteration 2701 / 9800) loss: 0.793451
(Iteration 2751 / 9800) loss: 1.217258
(Iteration 2801 / 9800) loss: 1.016398
(Iteration 2851 / 9800) loss: 0.805968
(Iteration 2901 / 9800) loss: 1.175808
(Epoch 6 / 20) train acc: 0.656000; val_acc: 0.605000
(Iteration 2951 / 9800) loss: 0.941986
(Iteration 3001 / 9800) loss: 0.982846
(Iteration 3051 / 9800) loss: 1.005875
(Iteration 3101 / 9800) loss: 0.929024
(Iteration 3151 / 9800) loss: 0.911446
(Iteration 3201 / 9800) loss: 0.831955
(Iteration 3251 / 9800) loss: 0.983365
(Iteration 3301 / 9800) loss: 1.084643
(Iteration 3351 / 9800) loss: 0.882134
(Iteration 3401 / 9800) loss: 0.907717
(Epoch 7 / 20) train acc: 0.693000; val_acc: 0.643000
(Iteration 3451 / 9800) loss: 0.926780
(Iteration 3501 / 9800) loss: 0.871370
(Iteration 3551 / 9800) loss: 0.916746
(Iteration 3601 / 9800) loss: 0.957859
(Iteration 3651 / 9800) loss: 0.945472
(Iteration 3701 / 9800) loss: 0.957690
(Iteration 3751 / 9800) loss: 0.890934
(Iteration 3801 / 9800) loss: 0.926908
(Iteration 3851 / 9800) loss: 1.042560
(Iteration 3901 / 9800) loss: 0.858834
(Epoch 8 / 20) train acc: 0.662000; val_acc: 0.633000
(Iteration 3951 / 9800) loss: 0.949428

(Iteration 4001 / 9800) loss: 1.022474
(Iteration 4051 / 9800) loss: 0.992603
(Iteration 4101 / 9800) loss: 1.047574
(Iteration 4151 / 9800) loss: 0.931832
(Iteration 4201 / 9800) loss: 0.772472
(Iteration 4251 / 9800) loss: 0.980237
(Iteration 4301 / 9800) loss: 1.027331
(Iteration 4351 / 9800) loss: 0.847622
(Iteration 4401 / 9800) loss: 0.800675
(Epoch 9 / 20) train acc: 0.737000; val_acc: 0.638000
(Iteration 4451 / 9800) loss: 0.834715
(Iteration 4501 / 9800) loss: 0.677432
(Iteration 4551 / 9800) loss: 1.080039
(Iteration 4601 / 9800) loss: 0.716728
(Iteration 4651 / 9800) loss: 0.806540
(Iteration 4701 / 9800) loss: 0.927516
(Iteration 4751 / 9800) loss: 0.875024
(Iteration 4801 / 9800) loss: 1.042306
(Iteration 4851 / 9800) loss: 0.882808
(Epoch 10 / 20) train acc: 0.721000; val_acc: 0.647000
(Iteration 4901 / 9800) loss: 0.892720
(Iteration 4951 / 9800) loss: 0.772937
(Iteration 5001 / 9800) loss: 0.810763
(Iteration 5051 / 9800) loss: 0.746560
(Iteration 5101 / 9800) loss: 0.696696
(Iteration 5151 / 9800) loss: 0.815378
(Iteration 5201 / 9800) loss: 0.859296
(Iteration 5251 / 9800) loss: 0.845369
(Iteration 5301 / 9800) loss: 0.803124
(Iteration 5351 / 9800) loss: 0.678393
(Epoch 11 / 20) train acc: 0.741000; val_acc: 0.662000
(Iteration 5401 / 9800) loss: 0.717994
(Iteration 5451 / 9800) loss: 0.730793
(Iteration 5501 / 9800) loss: 0.832582
(Iteration 5551 / 9800) loss: 0.773037
(Iteration 5601 / 9800) loss: 0.845416
(Iteration 5651 / 9800) loss: 0.783072
(Iteration 5701 / 9800) loss: 0.775187
(Iteration 5751 / 9800) loss: 0.882563
(Iteration 5801 / 9800) loss: 1.027854
(Iteration 5851 / 9800) loss: 0.830116
(Epoch 12 / 20) train acc: 0.759000; val_acc: 0.651000
(Iteration 5901 / 9800) loss: 0.705709
(Iteration 5951 / 9800) loss: 0.968117
(Iteration 6001 / 9800) loss: 0.736151
(Iteration 6051 / 9800) loss: 0.869076
(Iteration 6101 / 9800) loss: 0.766911
(Iteration 6151 / 9800) loss: 0.927024

(Iteration 6201 / 9800) loss: 0.662418
(Iteration 6251 / 9800) loss: 0.801865
(Iteration 6301 / 9800) loss: 0.654079
(Iteration 6351 / 9800) loss: 0.785943
(Epoch 13 / 20) train acc: 0.754000; val_acc: 0.648000
(Iteration 6401 / 9800) loss: 0.883470
(Iteration 6451 / 9800) loss: 0.720435
(Iteration 6501 / 9800) loss: 0.812503
(Iteration 6551 / 9800) loss: 0.775922
(Iteration 6601 / 9800) loss: 0.895231
(Iteration 6651 / 9800) loss: 0.766190
(Iteration 6701 / 9800) loss: 0.717640
(Iteration 6751 / 9800) loss: 0.738909
(Iteration 6801 / 9800) loss: 0.755856
(Iteration 6851 / 9800) loss: 0.754592
(Epoch 14 / 20) train acc: 0.751000; val_acc: 0.651000
(Iteration 6901 / 9800) loss: 0.691371
(Iteration 6951 / 9800) loss: 0.669393
(Iteration 7001 / 9800) loss: 0.748512
(Iteration 7051 / 9800) loss: 0.754870
(Iteration 7101 / 9800) loss: 0.611698
(Iteration 7151 / 9800) loss: 0.620703
(Iteration 7201 / 9800) loss: 0.804109
(Iteration 7251 / 9800) loss: 0.721399
(Iteration 7301 / 9800) loss: 0.729168
(Epoch 15 / 20) train acc: 0.797000; val_acc: 0.663000
(Iteration 7351 / 9800) loss: 0.794741
(Iteration 7401 / 9800) loss: 0.685225
(Iteration 7451 / 9800) loss: 0.659901
(Iteration 7501 / 9800) loss: 0.654125
(Iteration 7551 / 9800) loss: 0.897542
(Iteration 7601 / 9800) loss: 0.702215
(Iteration 7651 / 9800) loss: 0.958143
(Iteration 7701 / 9800) loss: 0.776002
(Iteration 7751 / 9800) loss: 0.690163
(Iteration 7801 / 9800) loss: 0.636967
(Epoch 16 / 20) train acc: 0.773000; val_acc: 0.651000
(Iteration 7851 / 9800) loss: 0.746292
(Iteration 7901 / 9800) loss: 0.577357
(Iteration 7951 / 9800) loss: 0.663366
(Iteration 8001 / 9800) loss: 0.851368
(Iteration 8051 / 9800) loss: 0.727580
(Iteration 8101 / 9800) loss: 0.683947
(Iteration 8151 / 9800) loss: 0.626854
(Iteration 8201 / 9800) loss: 0.800301
(Iteration 8251 / 9800) loss: 0.736672
(Iteration 8301 / 9800) loss: 0.818214
(Epoch 17 / 20) train acc: 0.760000; val_acc: 0.650000

(Iteration 8351 / 9800) loss: 0.802396
(Iteration 8401 / 9800) loss: 0.771348
(Iteration 8451 / 9800) loss: 0.622872
(Iteration 8501 / 9800) loss: 0.513887
(Iteration 8551 / 9800) loss: 0.539145
(Iteration 8601 / 9800) loss: 0.769144
(Iteration 8651 / 9800) loss: 0.708705
(Iteration 8701 / 9800) loss: 0.758026
(Iteration 8751 / 9800) loss: 0.806010
(Iteration 8801 / 9800) loss: 0.568705
(Epoch 18 / 20) train acc: 0.800000; val_acc: 0.661000
(Iteration 8851 / 9800) loss: 0.646179
(Iteration 8901 / 9800) loss: 0.531616
(Iteration 8951 / 9800) loss: 0.604449
(Iteration 9001 / 9800) loss: 0.779766
(Iteration 9051 / 9800) loss: 0.622247
(Iteration 9101 / 9800) loss: 0.680818
(Iteration 9151 / 9800) loss: 0.428375
(Iteration 9201 / 9800) loss: 0.928653
(Iteration 9251 / 9800) loss: 0.379808
(Iteration 9301 / 9800) loss: 0.789477
(Epoch 19 / 20) train acc: 0.817000; val_acc: 0.646000
(Iteration 9351 / 9800) loss: 0.763423
(Iteration 9401 / 9800) loss: 0.797257
(Iteration 9451 / 9800) loss: 0.708895
(Iteration 9501 / 9800) loss: 0.808784
(Iteration 9551 / 9800) loss: 0.568338
(Iteration 9601 / 9800) loss: 0.553148
(Iteration 9651 / 9800) loss: 0.637789
(Iteration 9701 / 9800) loss: 0.642199
(Iteration 9751 / 9800) loss: 0.614586
(Epoch 20 / 20) train acc: 0.807000; val_acc: 0.652000


```

1 import numpy as np
2
3 from nndl.layers import *
4 from nndl.conv_layers import *
5 from cs231n.fast_layers import *
6 from nndl.layer_utils import *
7 from nndl.conv_layer_utils import *
8
9 import pdb
10
11 """
12 This code was originally written for CS 231n at Stanford University
13 (cs231n.stanford.edu). It has been modified in various areas for use in the
14 ECE 239AS class at UCLA. This includes the descriptions of what code to
15 implement as well as some slight potential changes in variable names to be
16 consistent with class nomenclature. We thank Justin Johnson & Serena Yeung
17 for
18 permission to use this code. To see the original version, please visit
19 cs231n.stanford.edu.
20 """
21 class ThreeLayerConvNet(object):
22     """
23     A three-layer convolutional network with the following architecture:
24
25     conv - relu - 2x2 max pool - affine - relu - affine - softmax
26
27     The network operates on minibatches of data that have shape (N, C, H, W)
28     consisting of N images, each with height H and width W and with C input
29     channels.
30     """
31
32     def __init__(self, input_dim=(3, 32, 32), num_filters=32, filter_size=7,
33                 hidden_dim=100, num_classes=10, weight_scale=1e-3, reg=0.0,
34                 dtype=np.float32, use_batchnorm=False):
35         """
36         Initialize a new network.
37
38         Inputs:
39         - input_dim: Tuple (C, H, W) giving size of input data
40         - num_filters: Number of filters to use in the convolutional layer
41         - filter_size: Size of filters to use in the convolutional layer
42         - hidden_dim: Number of units to use in the fully-connected hidden layer
43         - num_classes: Number of scores to produce from the final affine layer.
44         - weight_scale: Scalar giving standard deviation for random
45         initialization
46           of weights.
47         - reg: Scalar giving L2 regularization strength
48         - dtype: numpy datatype to use for computation.
49         """
50         self.use_batchnorm = use_batchnorm
51         self.params = {}
52         self.reg = reg
53         self.dtype = dtype
54
55         # ===== #
56         # YOUR CODE HERE:
57         # Initialize the weights and biases of a three layer CNN. To
58         initialize:

```

```

58 # - the biases should be initialized to zeros.
59 # - the weights should be initialized to a matrix with entries
60 #     drawn from a Gaussian distribution with zero mean and
61 #     standard deviation given by weight_scale.
62 # ===== #
63
64 C, H, W = input_dim
65 pad = (filter_size - 1) // 2
66 conv_output_dims = (1 + (H + 2 * pad - filter_size))
67 pool_output_dims = int((conv_output_dims - 2) / 2 + 1) * 2 * num_filters
68
69 self.params['W1'] = weight_scale * np.random.randn(num_filters, C,
filter_size, filter_size)
70 self.params['b1'] = np.zeros(num_filters)
71 self.params['W2'] = weight_scale * np.random.randn(pool_output_dims,
hidden_dim)
72 self.params['b2'] = np.zeros(hidden_dim)
73 self.params['W3'] = weight_scale * np.random.randn(hidden_dim,
num_classes)
74 self.params['b3'] = np.zeros(num_classes)
75
76 # ===== #
77 # END YOUR CODE HERE
78 # ===== #
79
80 for k, v in self.params.items():
81     self.params[k] = v.astype(dtype)
82
83
84 def loss(self, X, y=None):
85     """
86     Evaluate loss and gradient for the three-layer convolutional network.
87
88     Input / output: Same API as TwoLayerNet in fc_net.py.
89     """
90     W1, b1 = self.params['W1'], self.params['b1']
91     W2, b2 = self.params['W2'], self.params['b2']
92     W3, b3 = self.params['W3'], self.params['b3']
93
94     # pass conv_param to the forward pass for the convolutional layer
95     filter_size = W1.shape[2]
96     conv_param = {'stride': 1, 'pad': (filter_size - 1) / 2}
97
98     # pass pool_param to the forward pass for the max-pooling layer
99     pool_param = {'pool_height': 2, 'pool_width': 2, 'stride': 2}
100
101     scores = None
102
103     # ===== #
104     # YOUR CODE HERE:
105     # Implement the forward pass of the three layer CNN. Store the output
106     # scores as the variable "scores".
107     # ===== #
108
109     caches = []
110     out, cache = conv_relu_pool_forward(X, W1, b1, conv_param, pool_param)
111     conv_shape = out.shape # used for backward pass
112     caches.append(cache)
113     out = np.reshape(out, (out.shape[0], -1))
114     out, cache = affine_relu_forward(out, W2, b2)

```

```

115     caches.append(cache)
116     scores, cache = affine_forward(out, W3, b3)
117     caches.append(cache)
118
119     # ===== #
120     # END YOUR CODE HERE
121     # ===== #
122
123     if y is None:
124         return scores
125
126     loss, grads = 0, {}
127     # ===== #
128     # YOUR CODE HERE:
129     # Implement the backward pass of the three layer CNN. Store the grads
130     # in the grads dictionary, exactly as before (i.e., the gradient of
131     # self.params[k] will be grads[k]). Store the loss as "loss", and
132     # don't forget to add regularization on ALL weight matrices.
133     # ===== #
134
135     sm_loss, dout = softmax_loss(scores, y)
136     reg_loss = self.reg * 0.5 * (np.sum(self.params['W1']**2) +
np.sum(self.params['W2']**2) + np.sum(self.params['W3']**2))
137     loss = sm_loss + reg_loss
138
139     dout, grads['W3'], grads['b3'] = affine_backward(dout, caches.pop())
140     dout, grads['W2'], grads['b2'] = affine_relu_backward(dout, caches.pop())
141     dout = dout.reshape(conv_shape)
142     dout, grads['W1'], grads['b1'] = conv_relu_pool_backward(dout,
caches.pop())
143
144     for weights in ['W3', 'W2', 'W1']:
145         grads[weights] += self.reg * self.params[weights]
146
147     # ===== #
148     # END YOUR CODE HERE
149     # ===== #
150
151     return loss, grads
152
153 pass
154
155

```

```

1 import numpy as np
2 from nn.layers import *
3 import pdb
4
5 """
6 This code was originally written for CS 231n at Stanford University
7 (cs231n.stanford.edu). It has been modified in various areas for use in the
8 ECE 239AS class at UCLA. This includes the descriptions of what code to
9 implement as well as some slight potential changes in variable names to be
10 consistent with class nomenclature. We thank Justin Johnson & Serena Yeung
11 for
12 permission to use this code. To see the original version, please visit
13 cs231n.stanford.edu.
14 """
15 def conv_forward_naive(x, w, b, conv_param):
16     """
17     A naive implementation of the forward pass for a convolutional layer.
18
19     The input consists of N data points, each with C channels, height H and
20 width
21 W. We convolve each input with F different filters, where each filter spans
22 all C channels and has height HH and width WW.
23
24 Input:
25 - x: Input data of shape (N, C, H, W)
26 - w: Filter weights of shape (F, C, HH, WW)
27 - b: Biases, of shape (F,)
28 - conv_param: A dictionary with the following keys:
29     - 'stride': The number of pixels between adjacent receptive fields in the
30       horizontal and vertical directions.
31     - 'pad': The number of pixels that will be used to zero-pad the input.
32
33 Returns a tuple of:
34 - out: Output data, of shape (N, F, H', W') where H' and W' are given by
35      $H' = 1 + (H + 2 * pad - HH) / stride$ 
36      $W' = 1 + (W + 2 * pad - WW) / stride$ 
37 - cache: (x, w, b, conv_param)
38 """
39 out = None
40 pad = conv_param['pad']
41 stride = conv_param['stride']
42
43 # ===== #
44 # YOUR CODE HERE:
45 # Implement the forward pass of a convolutional neural network.
46 # Store the output as 'out'.
47 # Hint: to pad the array, you can use the function np.pad.
48 # ===== #
49 xpad = np.pad(x, ((0, 0), (0, 0), (pad, pad), (pad, pad)), 'constant',
50 constant_values=0)
51 N, _, H, W = xpad.shape
52 F, _, HH, WW = w.shape
53 H_new = int(1 + (H - HH) / stride)
54 W_new = int(1 + (W - WW) / stride)
55 out = np.zeros((N, F, H_new, W_new))
56 for n in range(N):
57     for f in range(F):
58         for i in range(0, H_new):

```

```

58         for j in range(0, W_new):
59             i_start = i * stride
60             j_start = j * stride
61             x_patch = xpad[n, :, i_start:i_start+HH, j_start:j_start+WW]
62             conv = np.sum(np.multiply(x_patch, w[f])) + b[f]
63             out[n, f, i, j] = conv
64
65         # ===== #
66         # END YOUR CODE HERE
67         # ===== #
68
69     cache = (x, w, b, conv_param)
70     return out, cache
71
72
73 def conv_backward_naive(dout, cache):
74     """
75     A naive implementation of the backward pass for a convolutional layer.
76
77     Inputs:
78     - dout: Upstream derivatives.
79     - cache: A tuple of (x, w, b, conv_param) as in conv_forward_naive
80
81     Returns a tuple of:
82     - dx: Gradient with respect to x
83     - dw: Gradient with respect to w
84     - db: Gradient with respect to b
85     """
86     dx, dw, db = None, None, None
87
88     N, F, out_height, out_width = dout.shape
89     x, w, b, conv_param = cache
90
91     stride, pad = [conv_param['stride'], conv_param['pad']]
92     xpad = np.pad(x, ((0,0), (0,0), (pad,pad), (pad,pad)), mode='constant')
93     num_filts, _, f_height, f_width = w.shape
94
95     # ===== #
96     # YOUR CODE HERE:
97     # Implement the backward pass of a convolutional neural network.
98     # Calculate the gradients: dx, dw, and db.
99     # ===== #
100
101     dx = np.zeros_like(xpad)
102     dw = np.zeros_like(w)
103     db = np.zeros_like(b)
104     for n in range(N):
105         for f in range(F):
106             db[f] += np.sum(dout[n, f])
107             for i in range(out_height):
108                 for j in range(out_width):
109                     i_start = i * stride
110                     j_start = j * stride
111                     i_end = i_start + f_height
112                     j_end = j_start + f_width
113                     x_patch = xpad[n, :, i_start:i_end, j_start:j_end]
114
115                     dx[n, :, i_start:i_end, j_start:j_end] += w[f] * dout[n, f, i, j]
116                     dw[f] += dout[n, f, i, j] * x_patch
117

```

```

118 dx = dx[:, :, pad:pad+x.shape[2], pad:pad+x.shape[3]]
119
120 # ===== #
121 # END YOUR CODE HERE
122 # ===== #
123
124 return dx, dw, db
125
126
127 def max_pool_forward_naive(x, pool_param):
128     """
129     A naive implementation of the forward pass for a max pooling layer.
130
131     Inputs:
132     - x: Input data, of shape (N, C, H, W)
133     - pool_param: dictionary with the following keys:
134       - 'pool_height': The height of each pooling region
135       - 'pool_width': The width of each pooling region
136       - 'stride': The distance between adjacent pooling regions
137
138     Returns a tuple of:
139     - out: Output data
140     - cache: (x, pool_param)
141     """
142     out = None
143
144     # ===== #
145     # YOUR CODE HERE:
146     #   Implement the max pooling forward pass.
147     # ===== #
148
149     N, C, H, W = x.shape
150     pool_h = pool_param['pool_height']
151     pool_w = pool_param['pool_width']
152     stride = pool_param['stride']
153     H_new = int(1 + (H-pool_h) / stride)
154     W_new = int(1 + (W-pool_w) / stride)
155     out = np.zeros((N, C, H_new, W_new))
156
157     for n in range(N):
158         for c in range(C):
159             for i in range(H_new):
160                 for j in range(W_new):
161                     i_start = i * stride
162                     j_start = j * stride
163                     i_end = i_start + pool_h
164                     j_end = j_start + pool_w
165                     out[n, c, i, j] = np.max(x[n, c, i_start:i_end, j_start:j_end])
166
167     # ===== #
168     # END YOUR CODE HERE
169     # ===== #
170     cache = (x, pool_param)
171     return out, cache
172
173 def max_pool_backward_naive(dout, cache):
174     """
175     A naive implementation of the backward pass for a max pooling layer.
176
177     Inputs:

```

```

178 - dout: Upstream derivatives
179 - cache: A tuple of (x, pool_param) as in the forward pass.
180
181 Returns:
182 - dx: Gradient with respect to x
183 """
184 dx = None
185 x, pool_param = cache
186 pool_height, pool_width, stride = pool_param['pool_height'],
pool_param['pool_width'], pool_param['stride']
187
188 # ===== #
189 # YOUR CODE HERE:
190 #   Implement the max pooling backward pass.
191 # ===== #
192 N, C, H, W = dout.shape
193
194 dx = np.zeros_like(x)
195
196 for n in range(N):
197     for c in range(C):
198         for i in range(H):
199             for j in range(W):
200                 i_start = i * stride
201                 j_start = j * stride
202                 i_end = i_start + pool_height
203                 j_end = j_start + pool_width
204                 x_patch = x[n, c, i_start:i_end, j_start:j_end]
205                 mask = x_patch == np.max(x_patch)
206                 dx[n, c, i_start:i_end, j_start:j_end] += mask * dout[n, c, i, j]
207
208 # ===== #
209 # END YOUR CODE HERE
210 # ===== #
211
212 return dx
213 def spatial_batchnorm_forward(x, gamma, beta, bn_param):
214     """
215     Computes the forward pass for spatial batch normalization.
216
217     Inputs:
218     - x: Input data of shape (N, C, H, W)
219     - gamma: Scale parameter, of shape (C,)
220     - beta: Shift parameter, of shape (C,)
221     - bn_param: Dictionary with the following keys:
222       - mode: 'train' or 'test'; required
223       - eps: Constant for numeric stability
224       - momentum: Constant for running mean / variance. momentum=0 means that
225         old information is discarded completely at every time step, while
226         momentum=1 means that new information is never incorporated. The
227         default of momentum=0.9 should work well in most situations.
228       - running_mean: Array of shape (D,) giving running mean of features
229       - running_var: Array of shape (D,) giving running variance of features
230
231     Returns a tuple of:
232     - out: Output data, of shape (N, C, H, W)
233     - cache: Values needed for the backward pass
234     """
235     out, cache = None, None
236

```

```

237 # ===== #
238 # YOUR CODE HERE:
239 #   Implement the spatial batchnorm forward pass.
240 #
241 #   You may find it useful to use the batchnorm forward pass you
242 #   implemented in HW #4.
243 # ===== #
244
245 N, C, H, W = x.shape
246 x = x.transpose(0, 2, 3, 1).reshape((-1, C))
247 out, cache = batchnorm_forward(x, gamma, beta, bn_param)
248 out = out.reshape((N, H, W, C)).transpose(0, 3, 1, 2)
249
250 # ===== #
251 # END YOUR CODE HERE
252 # ===== #
253
254 return out, cache
255
256 def spatial_batchnorm_backward(dout, cache):
257     """
258     Computes the backward pass for spatial batch normalization.
259
260     Inputs:
261     - dout: Upstream derivatives, of shape (N, C, H, W)
262     - cache: Values from the forward pass
263
264     Returns a tuple of:
265     - dx: Gradient with respect to inputs, of shape (N, C, H, W)
266     - dgamma: Gradient with respect to scale parameter, of shape (C,)
267     - dbeta: Gradient with respect to shift parameter, of shape (C,)
268     """
269     dx, dgamma, dbeta = None, None, None
270
271 # ===== #
272 # YOUR CODE HERE:
273 #   Implement the spatial batchnorm backward pass.
274 #
275 #   You may find it useful to use the batchnorm forward pass you
276 #   implemented in HW #4.
277 # ===== #
278
279 N, C, H, W = dout.shape
280 dout = dout.transpose(0, 2, 3, 1).reshape((-1, C))
281 dx, dgamma, dbeta = batchnorm_backward(dout, cache)
282 dx = dx.reshape((N, H, W, C)).transpose(0, 3, 1, 2)
283
284 # ===== #
285 # END YOUR CODE HERE
286 # ===== #
287
288 return dx, dgamma, dbeta

```